INFO 251: Applied ML

Decision Trees

Should I do the laundry? A decision tree:

Quiz 1 results

Average Score	High Score	Low Score	Standard Deviation	Average Time
89%	100%	25%	4.13	26:50

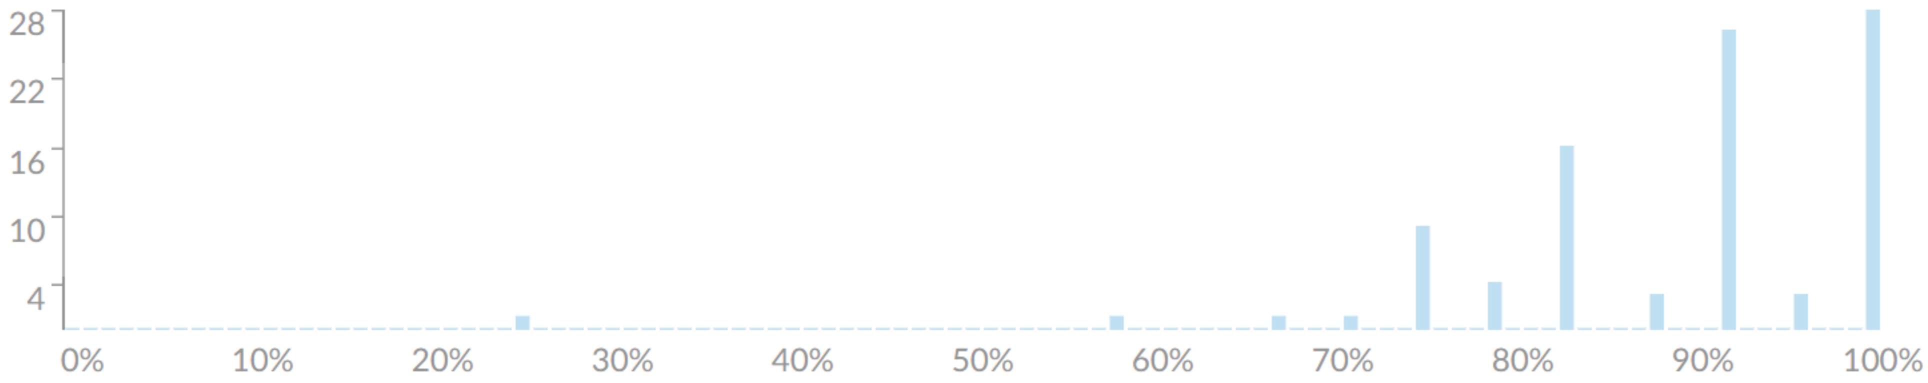

Quiz 1 results

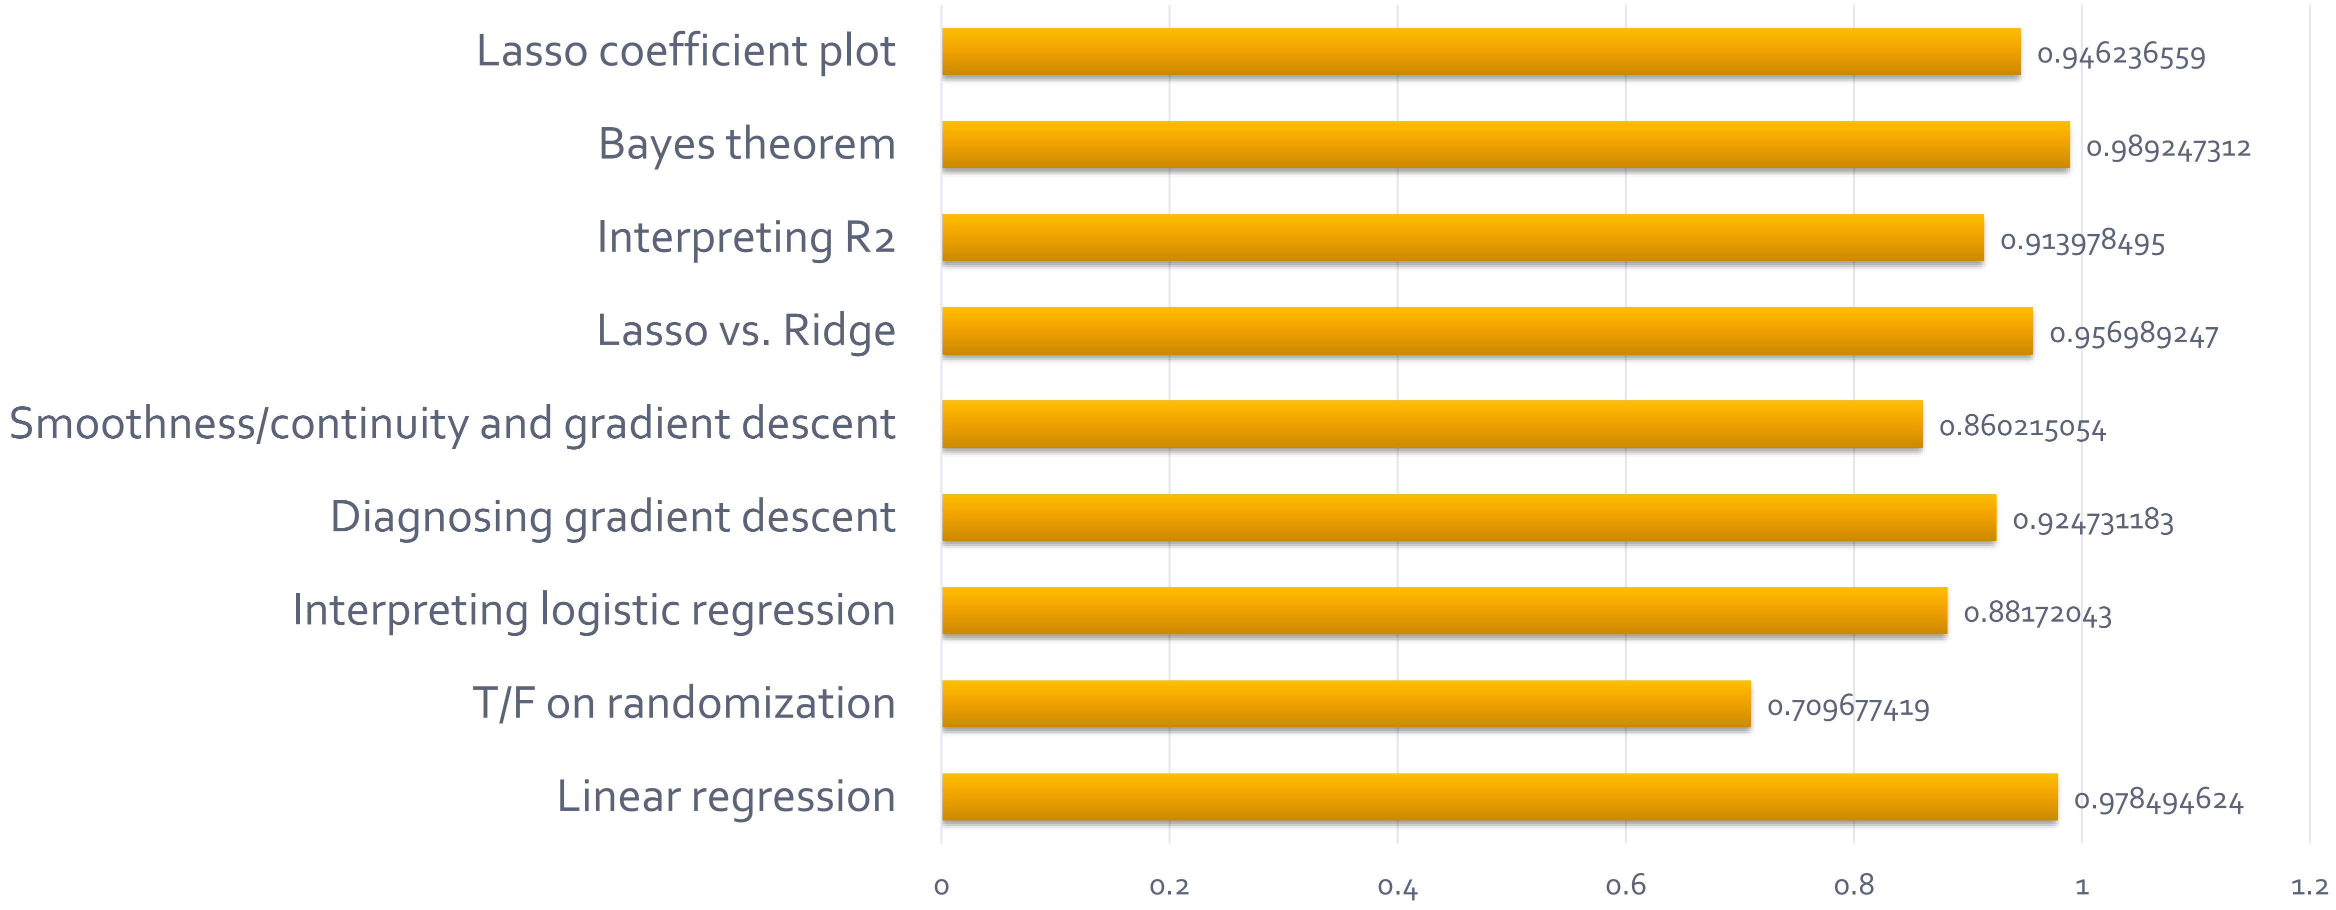

Quiz 1 Review

Attempts: 93 out of 93

+0.37

When treatm ... dex

treatment) ou

estimate of th

True
False

Randomization (Lecture 2)

- Randomize the treatment status
- Ensures that attributes (observable and unobservable) of treated and untreated individuals are the same, on average
- *Under randomized treatment, a simple difference between outcomes in treated and control units gives unbiased estimate of impact*

Quiz 1 Review

Attempts: 93 out of 93

What is the Double-Difference estimate of impact?

- Before intervention, Treatment group outcome = 150
- After intervention, Treatment group outcome = 144
- Before intervention, Control group outcome = 151
- After intervention, Control group outcome = 155

-10.00	81 respondents	87 %
Something Else	12 respondents	13 %

87% answered correctly

Quiz 1 Review

Attempts: 93 out of 93

If the cost function is continuous and differentiable, and the learning rate is sufficiently small, gradient descent is guaranteed to eventually converge to the global minimum.

True	13 respondents	14 %
False	80 respondents	**86 %**

+0.45

Discrimination Index

86% answered correctly

Quiz 1 Review

Attempts: 93 out of 93

When measuring the distance between two m-dimensional points x_i and x_j, a common distance metric is the L-norm, defined as:

$$D^n\left(x_i, x_j\right) = \sqrt[n]{\sum_{m=1}^{M}\left(x_{im} - x_{jm}\right)^n}$$

When $n = 0$, what is the distance between the points $x_1 = (0,0,0)$ and $x_2 = (2,3,6)$? i.e., what is $D^0\left((0,0,0),(2,3,6)\right)$? Please assume that $\sqrt[0]{Z} = Z$.

3.00	80 respondents	86 %
Something Else	13 respondents	14 %

86% answered correctly

Quiz 1 Review

Attempts: 93

Which of t
non-linear

Note: This
implement
a linear alg

- Linear Regr
- Logistic Re
- **k-Nearest**
- **k-Nearest**
- Lasso-regu
- None of th

Logistic Regression: Example

- Example: admission vs. GRE and GPA
  1. Start with raw data
  2. Fit logistic regression
  3. Threshold converts g(z) to classification

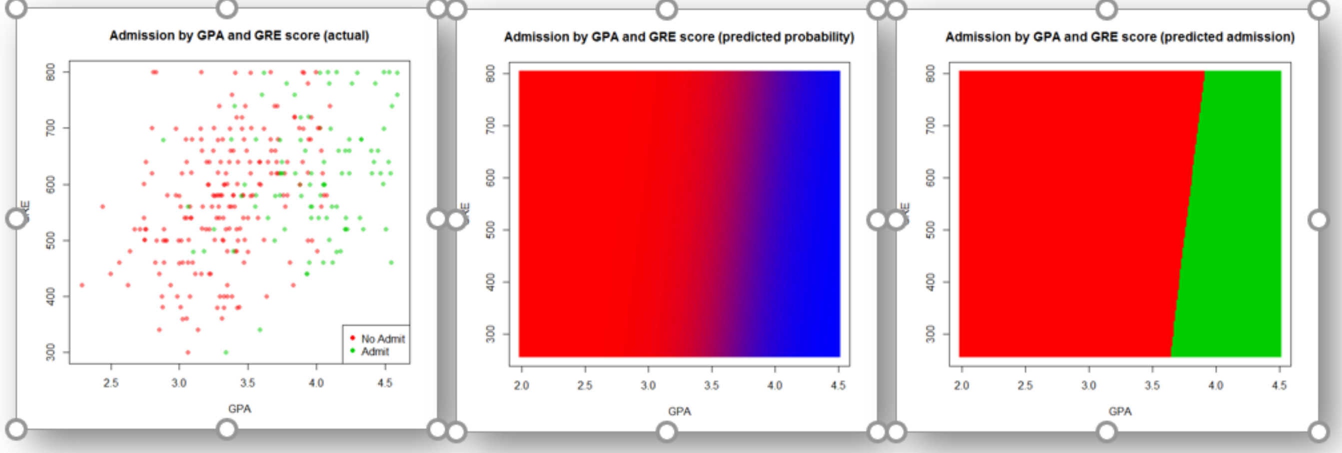

Course Outline

- Causal Inference and Research Design
  - Experimental methods
  - Non-experiment methods
- Machine Learning
  - Design of Machine Learning Experiments
  - Linear Models and Gradient Descent
  - **Non-linear models**
  - Neural models
  - Unsupervised Learning
  - Practicalities, Fairness, Bias
- Special topics

Outline

- Decision Trees
  - Introduction
  - Representation
  - Algorithms
  - Splitting
  - Extended example
  - Overfitting and Pruning
  - Extensions

Key Concepts (this lecture)

- Churn prediction
- Decision boundaries
- Hyper-rectangles
- Splitting
- Information gain
- Recursive tree building
- Overfitting trees
- Pruning trees

Decision Trees

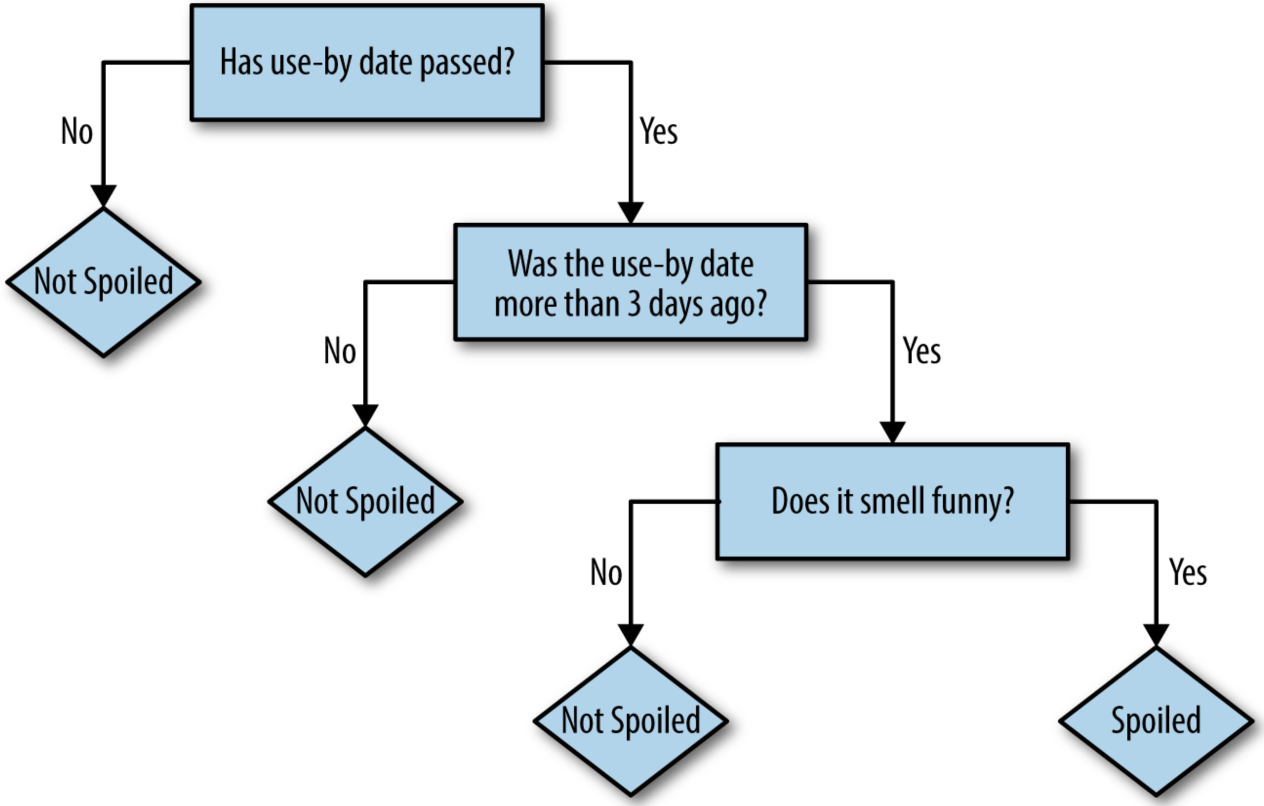

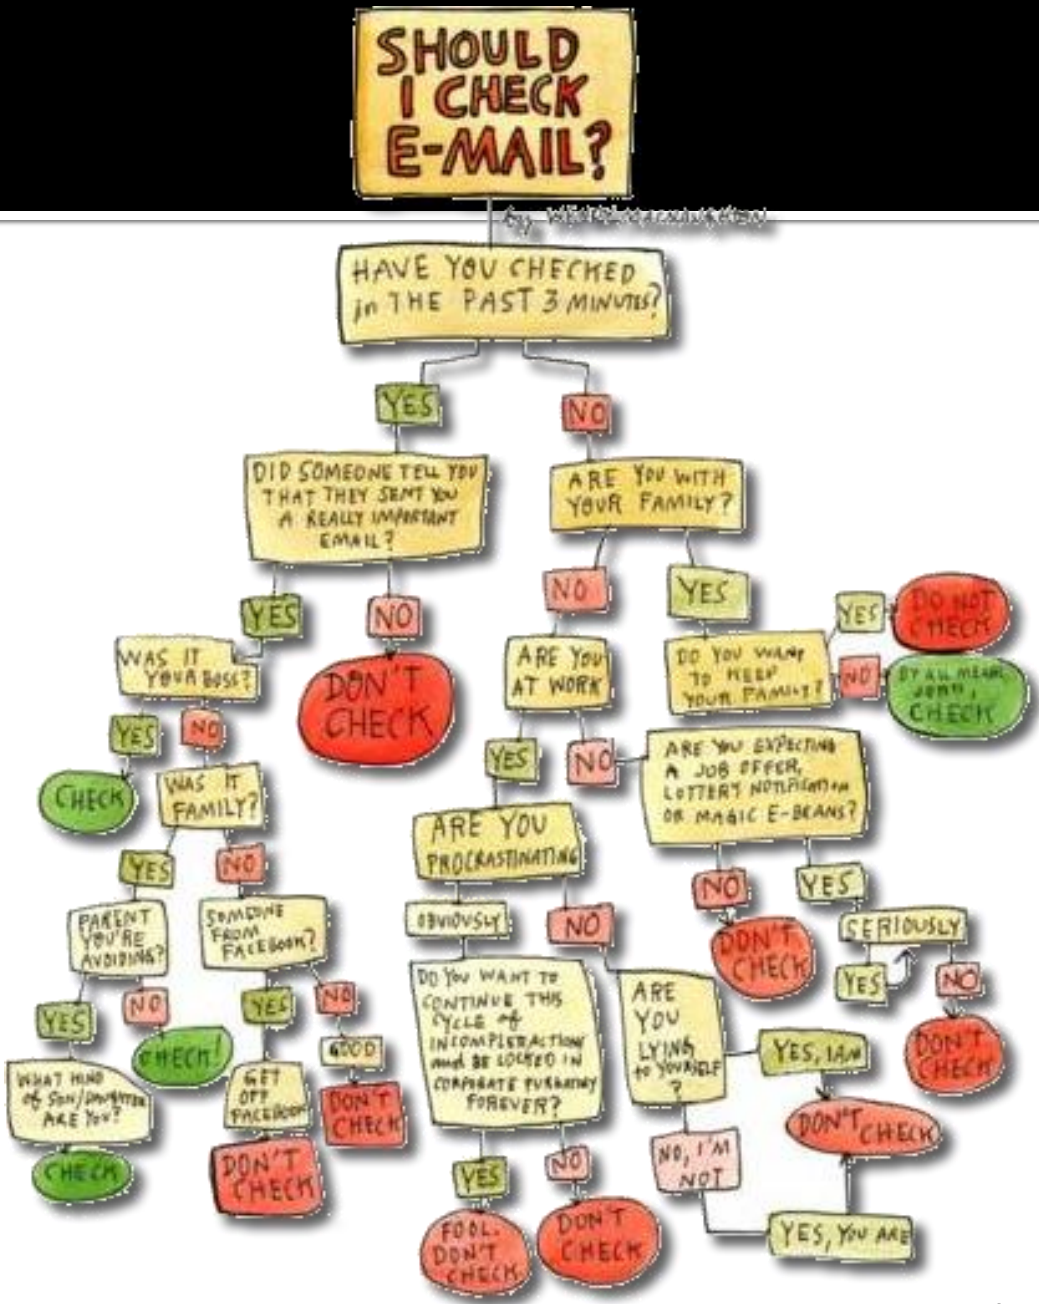

Example: Customer churn

- Goal: reduce customer churn
  - E.g., target customers to encourage retention
  - Start by predicting who is likely to churn
  - What features to include?

Variable	Explanation
COLLEGE	Is the customer college educated?
INCOME	Annual income
OVERAGE	Average overcharges per month
LEFTOVER	Average number of leftover minutes per month
HOUSE	Estimated value of dwelling (from census tract)
HANDSET_PRICE	Cost of phone
LONG_CALLS_PER_MONTH	Average number of long calls (15 mins or over) per month
AVERAGE_CALL_DURATION	Average duration of a call
REPORTED_SATISFACTION	Reported level of satisfaction
REPORTED_USAGE_LEVEL	Self-reported usage level
LEAVE	*Target variable: Did the customer stay or leave (churn)?*

[Provost and Fawcett]

Example: Customer churn

- A simple approach: "Forward Selection"
  - Add features to a model (e.g., kNN, logistic regression, Naïve Bayes) one at a time, based on the relevance of that feature
  - How to define "relevance"?
    - Unconditional correlation
    - Information gain (we'll define this soon)

Rank	Info. Gain	Attribute name
1	0.0461296	HOUSE
2	0.0435518	OVERAGE
3	0.0350337	LONG_CALLS_PER_MON
4	0.013648	LEFTOVER
5	0.0100534	INCOME
6	0.0088899	AVG_CALL_DURATION
7	0.007624	HANDSET_PRICE
8	0.0003062	SATISFACTION
9	0.0001553	COLLEGE
10	0.0000388	USAGE_LEVEL

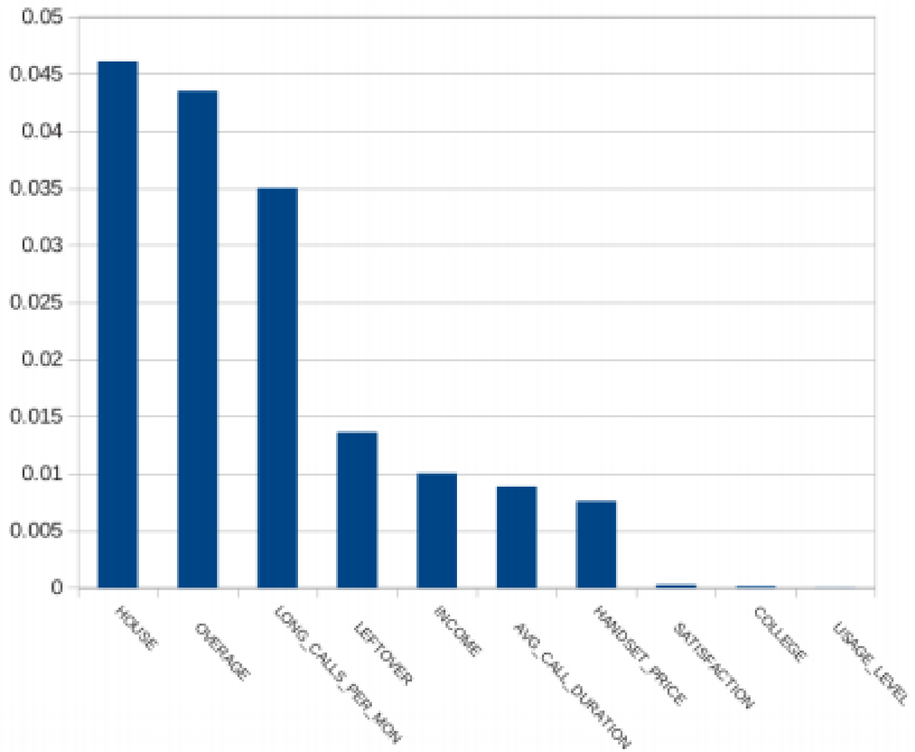

Example: Customer churn

- Decision Tree
  - Idea: Build tree from training data to explain labeled examples
  - Progressively add features that are most informative
  - Keep tree as small and as simple as possible (Occam's razor)
- Different from forward selection
  - Decisions are made conditional on existing tree
  - E.g., second variable not necessarily OVERAGE

Rank	Info. Gain	Attribute name
1	0.0461296	HOUSE
2	0.0435518	OVERAGE
3	0.0350337	LONG_CALLS_PER_MON
4	0.013648	LEFTOVER
5	0.0100534	INCOME
6	0.0088899	AVG_CALL_DURATION
7	0.007624	HANDSET_PRICE
8	0.0003062	SATISFACTION
9	0.0001553	COLLEGE
10	0.0000388	USAGE_LEVEL

Example: Customer churn

- The final tree (churn dataset)

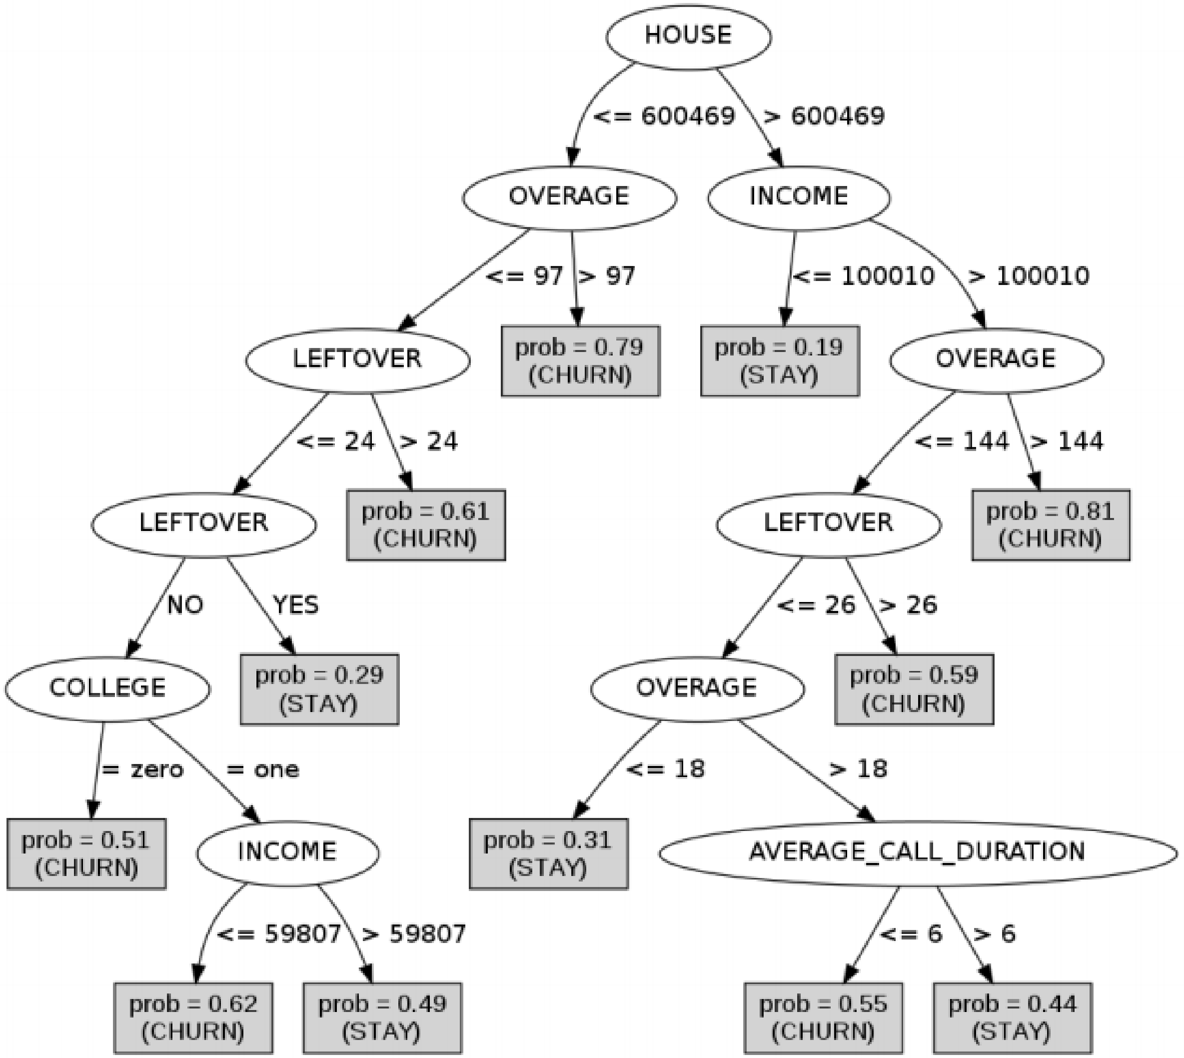

Decision trees

- Very popular, easy to interpret
- Reflects the logic of decision-making
- Arbitrarily complex (non-linear functions)
- Can be used for classification or regression
- Simple, fast algorithms for generating compact trees from data

Outline

- Decision Trees
  - Introduction
  - **Representation**
  - Algorithms
  - Splitting
  - Extended example
  - Overfitting and Pruning
  - Extensions

Representation

- Internal nodes test the value of a feature
  - Categorical
  - Binary
  - Continuous
- Branches indicate possible values for feature
- Leaf nodes output a predicted class or value

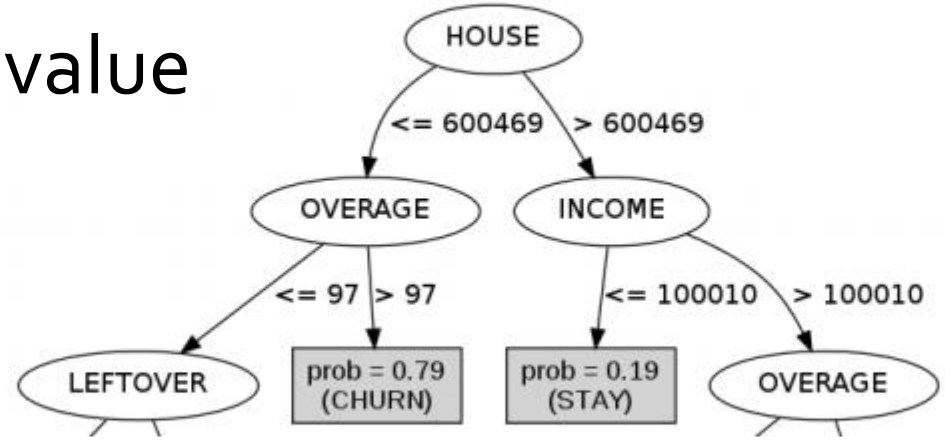

Simple Example: XOR

- Expressiveness
  - Any logical function that can be encoded in a truth table can be expressed in a decision tree!
  - Why? For any truth table, use each variables as a split

x_1	x_2	y
1	1	-1
1	0	1
0	1	1
0	0	-1

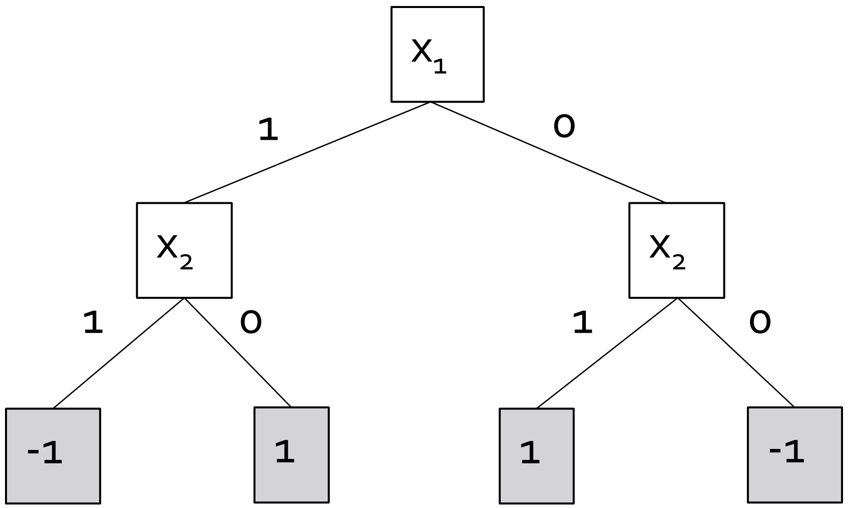

Decision boundaries

- What does the decision boundary of a decision tree look like?
  - Compare to k-NN? Logistic regression?

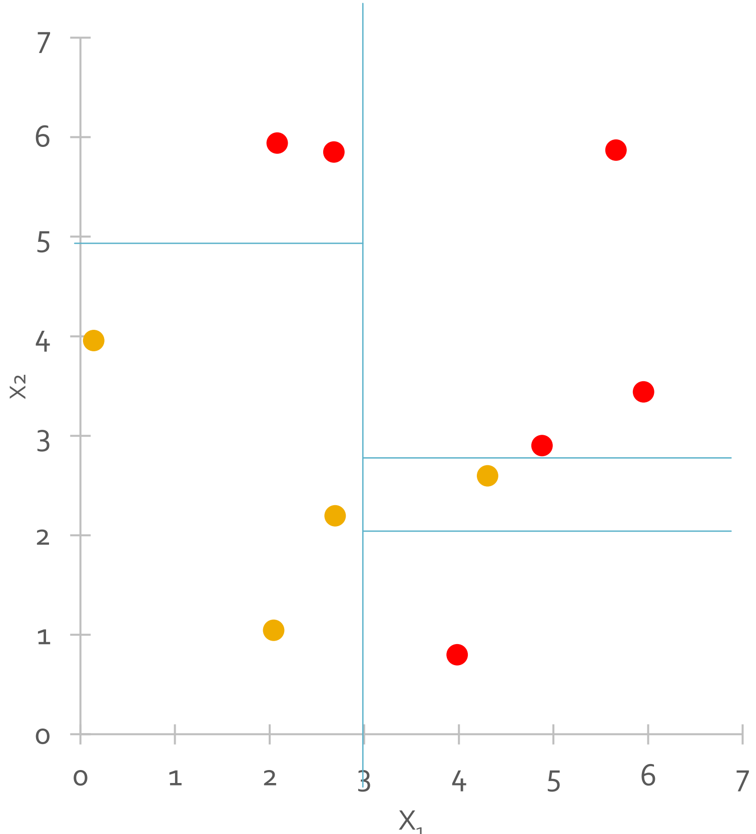

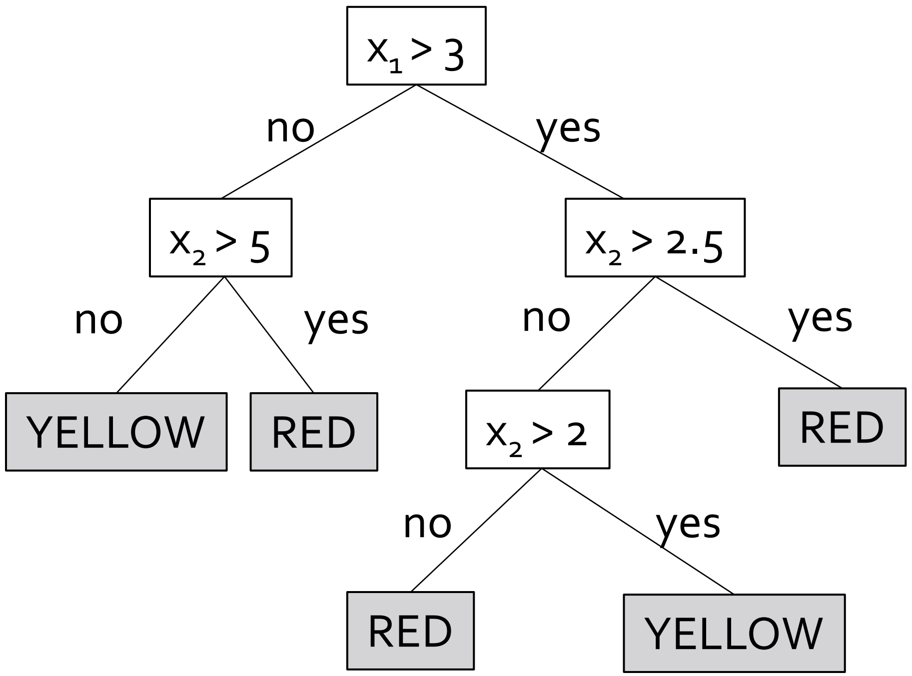

Decision boundaries

- "Hyper-rectangles" partition high-dimensional space
- Diagonal lines approximated as step function

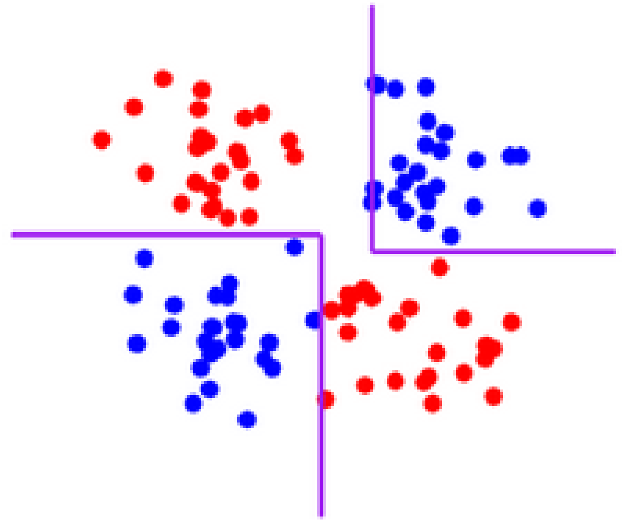

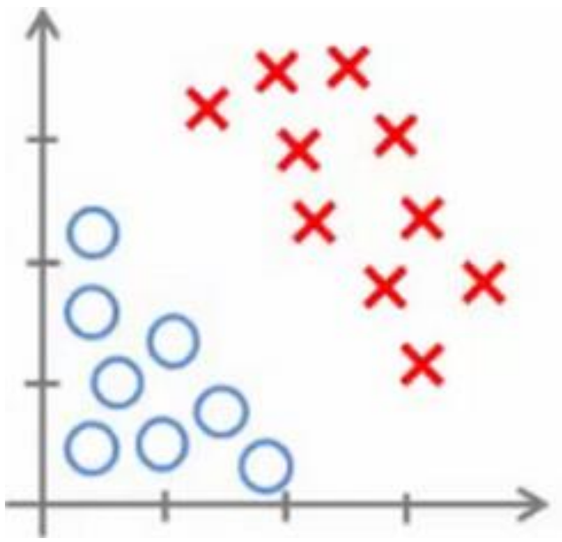

Hypothesis space

- How many possible decision trees over *N* binary variables?
  - = number of distinct truth tables with 2^N rows
  - = $2^\wedge(2^N)$
  - e.g. 6 attributes: 18,446,744,073,709,551,616 possible trees
  - Be careful!
- More expressive hypothesis spaces…
  - Increases chance that target function can be expressed
  - Increases hypotheses consistent with training data
  - Means we can get better predictions
  - But we may fail to generalize

Good vs. bad trees

- Many trees perfectly classify all training examples
  - A trivial model has a leaf for every training example
  - Doesn't matter what the root feature is
- But we want our tree to generalize
  - i.e., we want the smallest tree that explains the data

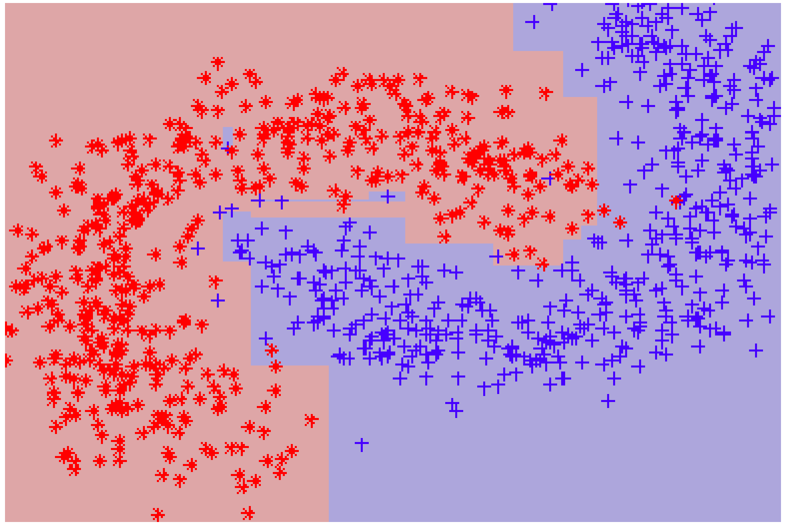

Intuition check

- True or False: A decision tree is able to recover non-linear decision boundary

Outline

- Decision Trees
  - Introduction
  - Representation
  - **Algorithms**
  - Splitting
  - Extended example
  - Overfitting and Pruning
  - Extensions

Building a tree (recursively)

- Goal: find a tree consistent with the training examples
- Strategy: (recursively) choose the most significant attribute as the root of the (sub)tree
- Example: binary features and binary labels

```
GrowTree(S):
      if y==0 for all <x,y> in S:
             return new leaf(0)
      else if y==1 for all <x,y> in S:
             return new leaf(1)
      else:
             choose best attribute x_j
             S0 = all <x,y> in S with x_j==0
             S1 = all <x,y> in S with x_j==1
             return new node(x_j, GrowTree(S0), GrowTree(S1))
```

Outline

- Decision Trees
  - Introduction
  - Representation
  - Algorithms
  - **Splitting**
  - Extended example
  - Overfitting and Pruning
  - Extensions

What is the “best” attribute?

- Good attributes split examples into pure subsets
- Should we split on COLLEGE or SATISFIED?
  - How to measure subset purity?

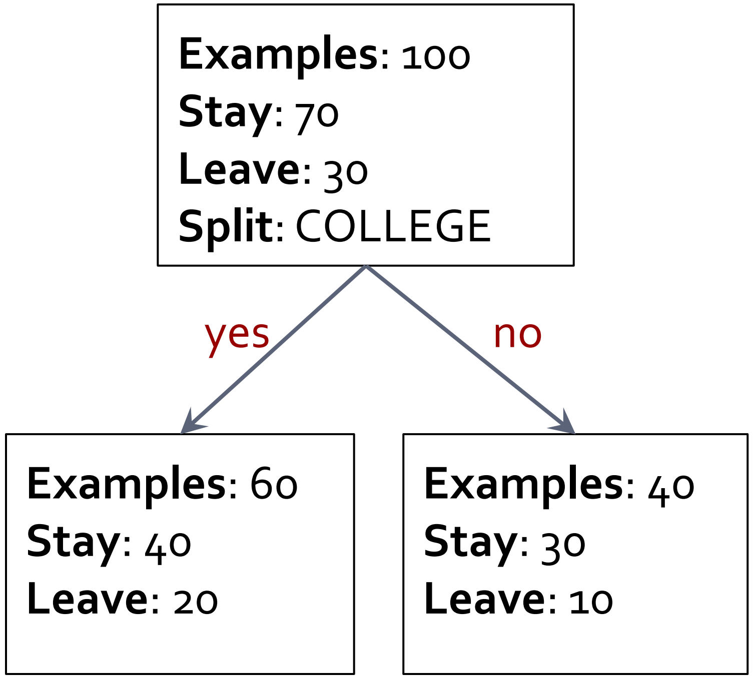

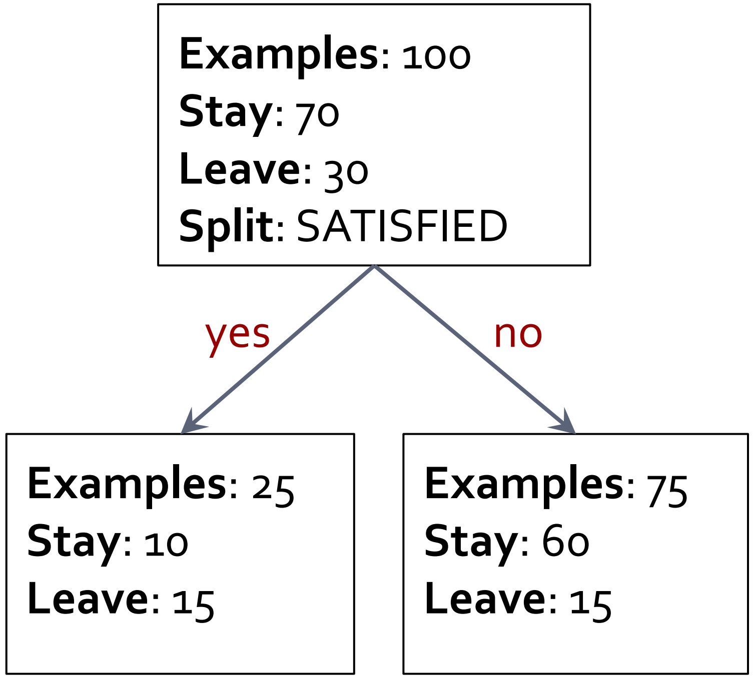

Information Theory

- Shannon's Game: predict the character

```
ARE WE THERE YE_
```

  - Explores mathematics of encoded messages
- How much information is conveyed by a single letter?
  - If the alphabet contains only one letter: 0 bits
  - With 27 equiprobable letters: 4.8 bits
  - Shannon's estimate of English characters: ~1 bit
- "Information" is the expected code length to convey a message

Claude Shannon
1916-2001

Entropy

- Entropy is a measure of uncertainty
  - More uncertainty requires longer codes
  - A distribution where one value has P(1) has no entropy
  - Uniform distribution maximizes entropy

Bits required

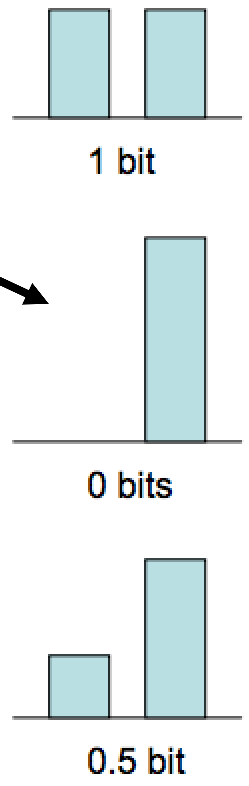

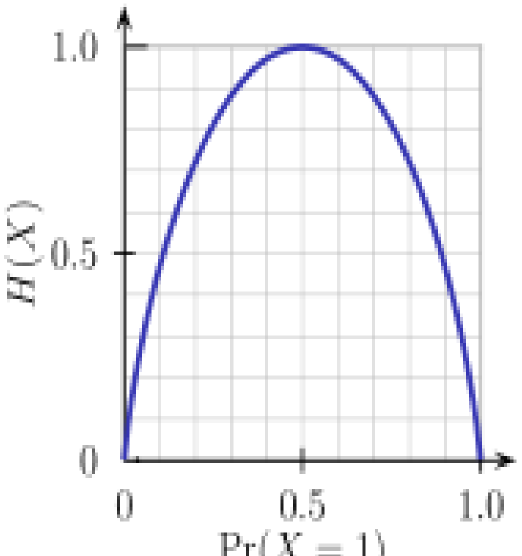

$$H = -\sum_i p_i(\log_2 p_i)$$

Information gain

- Information Gain: describes a change in entropy
  - Entropy before - Entropy after
- Example with positive information gain:
- Our example is trickier
  - Requires a calculator

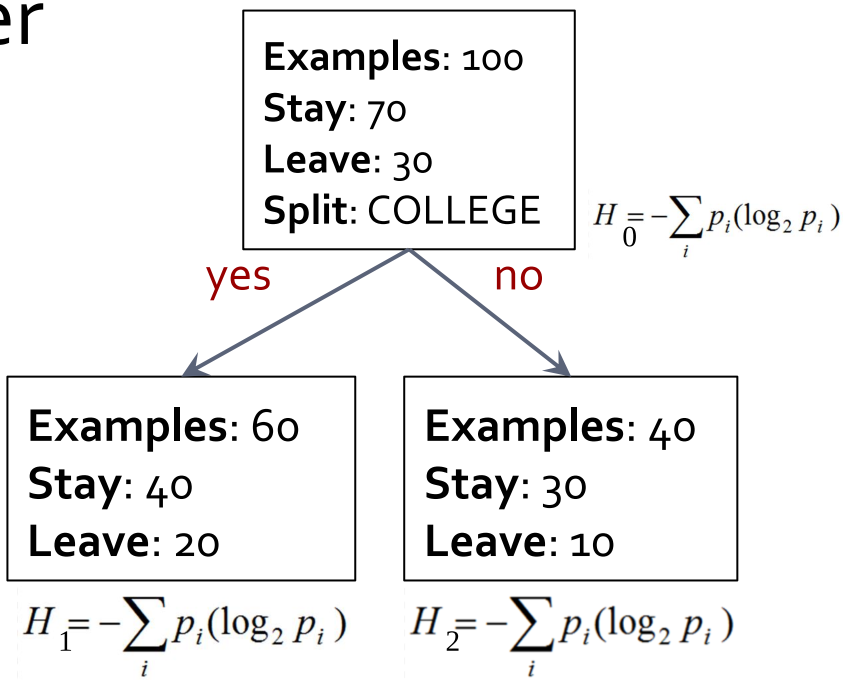

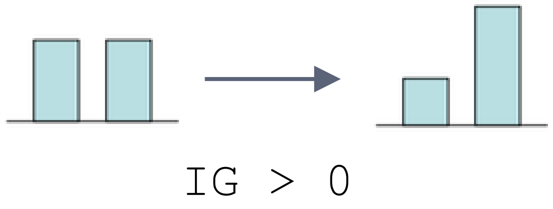

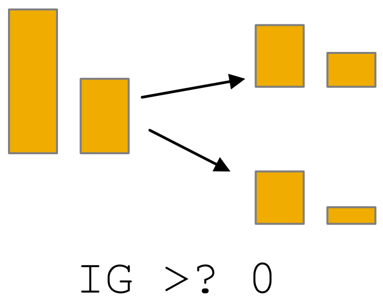

Information gain

- Before we calculate information gain, notice that the split produces branches of different sizes (60 vs. 40)
  - Solution: weight by the number of examples

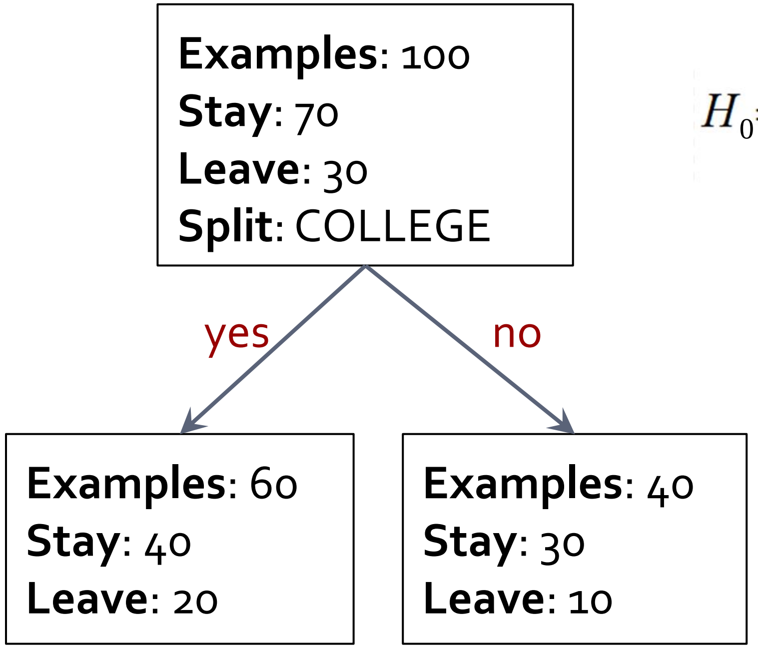

$$H_0 = -\sum_i p_i(\log_2 p_i)$$

$$H_1 = -\sum_i p_i(\log_2 p_i)$$

$$H_2 = -\sum_i p_i(\log_2 p_i)$$

Information gain: example

$$H = -\sum_i p_i(\log_2 p_i)$$

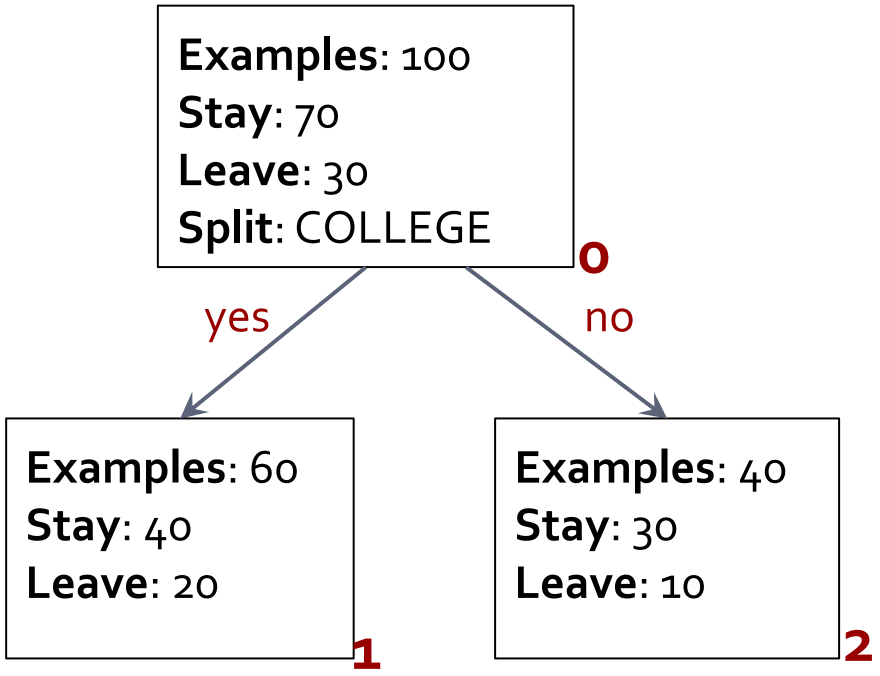

$$H_0 = -.7 \log(.7) - .3 \log(.3) = \mathbf{.88}$$

$$H_1 = -.66 \log(.66) - .33 \log(.33) = \mathbf{.92}$$

$$H_2 = -.75 \log(.75) - .25 \log(.25) = \mathbf{.81}$$

$$\begin{aligned} IG(COLLEGE) &= .88 - [.6(.92) + .4(.81)] \\ &= .88 - .87 \\ &= \mathbf{.01} \end{aligned}$$

(= **small gain**)

Outline

- Decision Trees
  - Introduction
  - Representation
  - Algorithms
  - Splitting
  - **Extended example**
  - Overfitting and Pruning
  - Extensions

Extended Example: Predicting MPG

mpg	cylinders	displacement	horsepower	weight	acceleration	modelyear	maker
good	4	low	low	low	high	75to78	asia
bad	6	medium	medium	medium	medium	70to74	america
bad	4	medium	medium	medium	low	75to78	europe
bad	8	high	high	high	low	70to74	america
bad	6	medium	medium	medium	medium	70to74	america
bad	4	low	medium	low	medium	70to74	asia
bad	4	low	medium	low	low	70to74	asia
bad	8	high	high	high	low	75to78	america
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
bad	8	high	high	high	low	70to74	america
good	8	high	medium	high	high	79to83	america
bad	8	high	high	high	low	75to78	america
good	4	low	low	low	low	79to83	america
bad	6	medium	medium	medium	high	75to78	america
good	4	medium	low	low	low	79to83	america
good	4	low	low	medium	high	79to83	america
bad	8	high	high	high	low	70to74	america
good	4	low	medium	low	medium	75to78	europe
bad	5	medium	medium	medium	medium	75to78	europe

[Andrew Moore]

The first split

- Each attribute is correlated with the target
- Calculate information gain for each possible split
  - (Note that attributes don't have to be binary)
- Choose the split that maximizes information gain

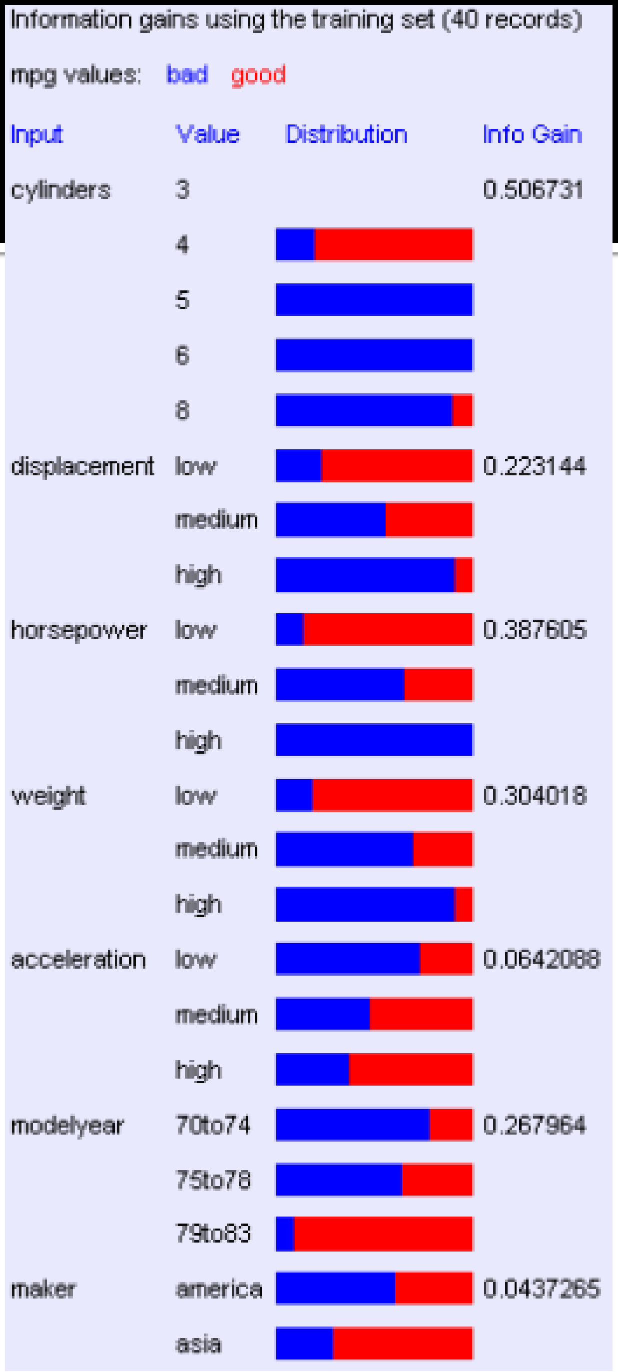

Information gains using the training set (40 records)

mpg values: bad good

Input	Value	Distribution	Info Gain
cylinders	3		0.506731
	4		
	5		
	6		
	8		
displacement	low		0.223144
	medium		
	high		
horsepower	low		0.387605
	medium		
	high		
weight	low		0.304018
	medium		
	high		
acceleration	low		0.0642088
	medium		
	high		
modelyear	70to74		0.267964
	75to78		
	79to83		
maker	america		0.0437265
	asia		

Depth-1 tree: “decision stump”

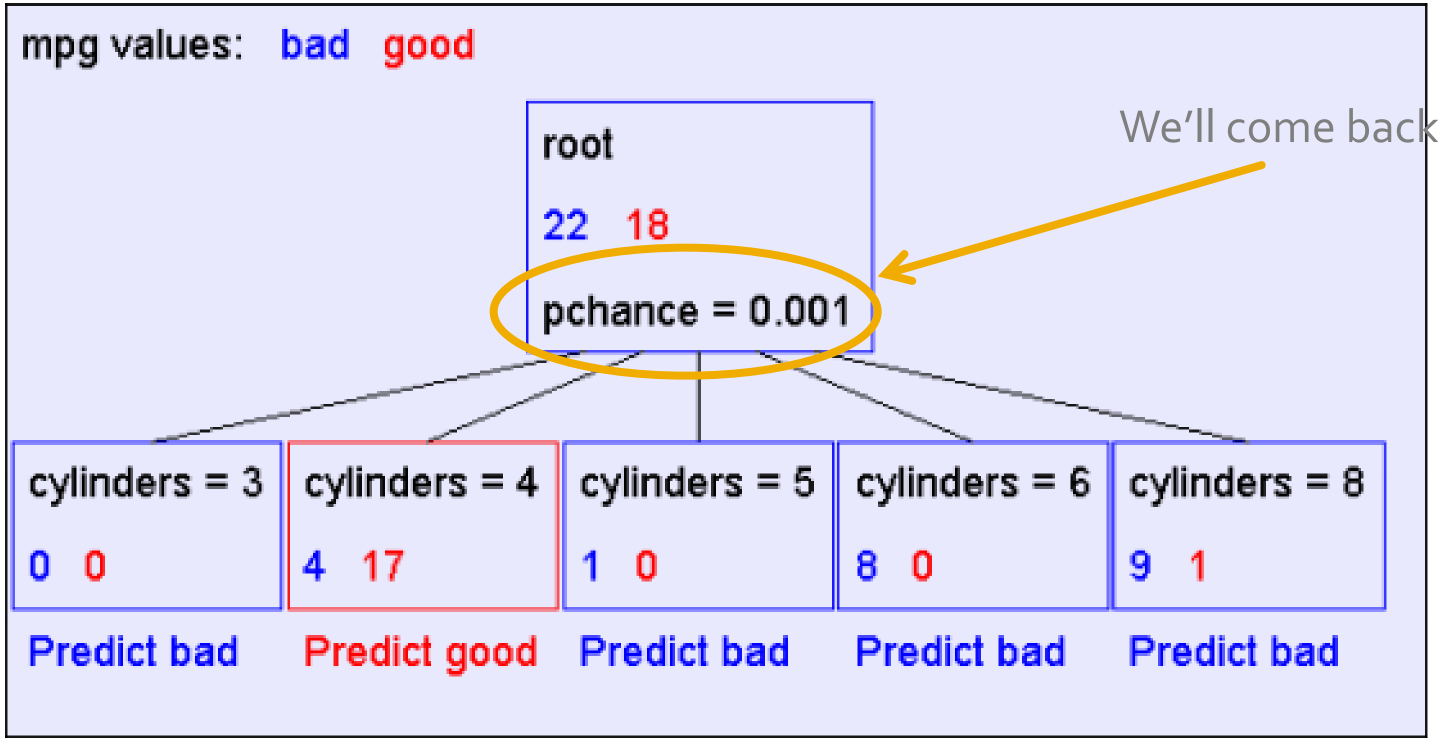

We’ll come back to this

Recursive tree building

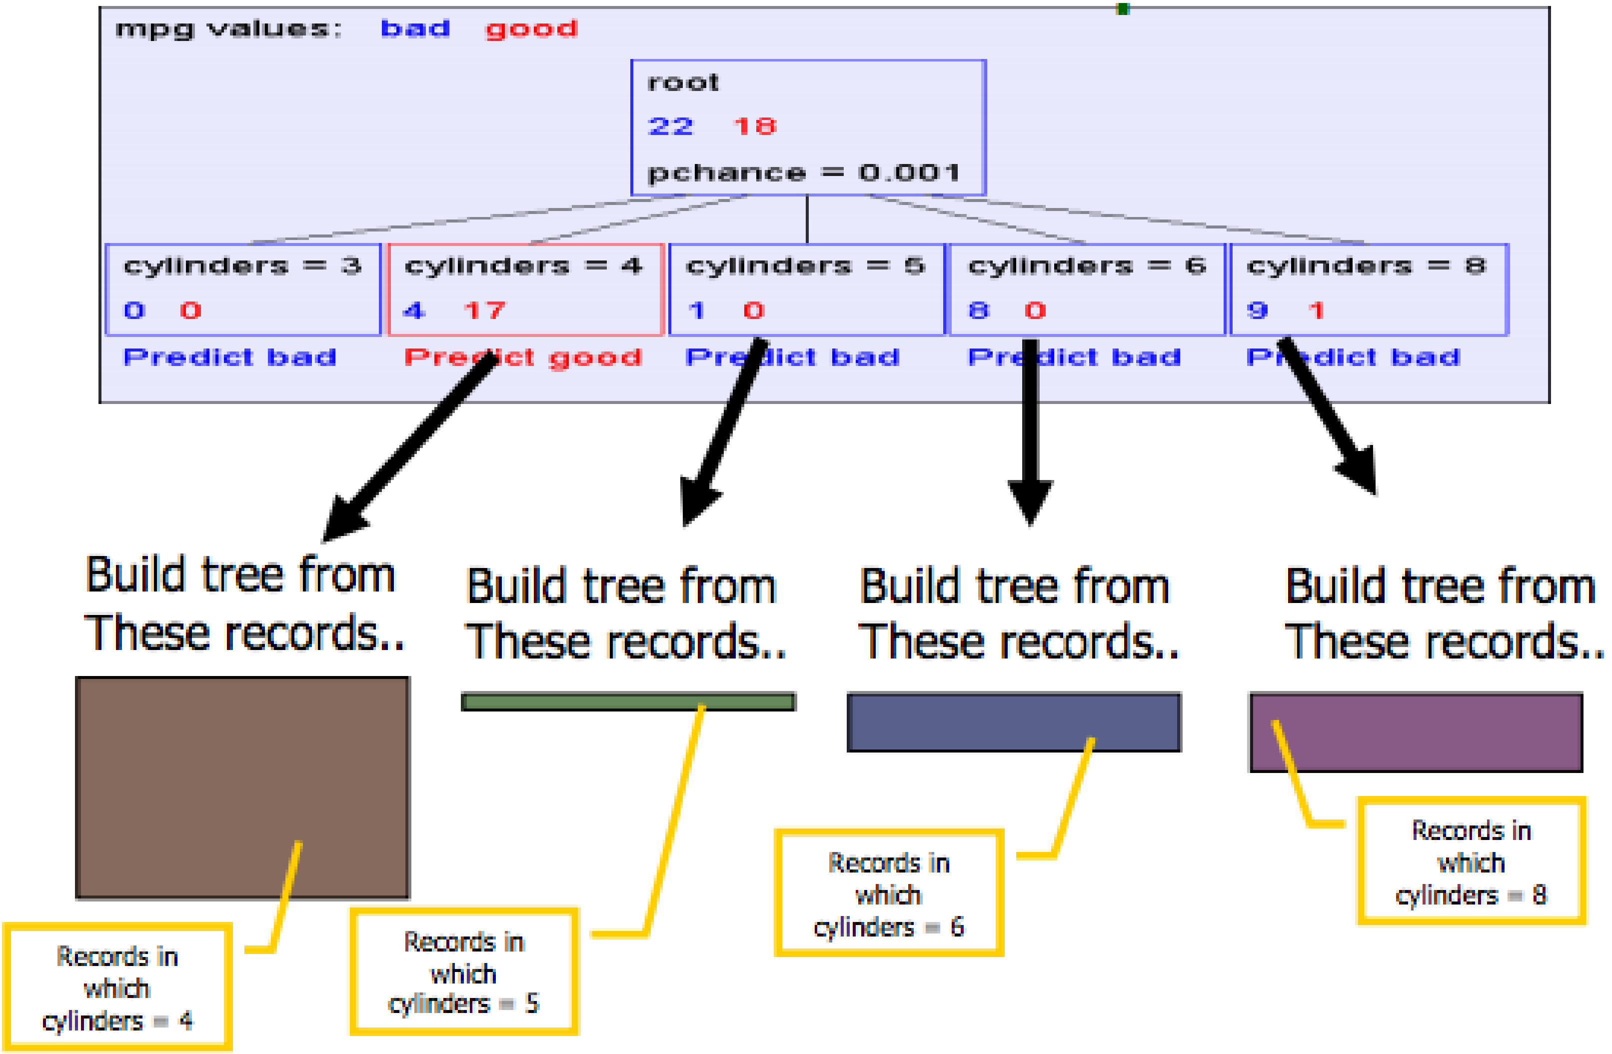

Second level (depth-2 tree)

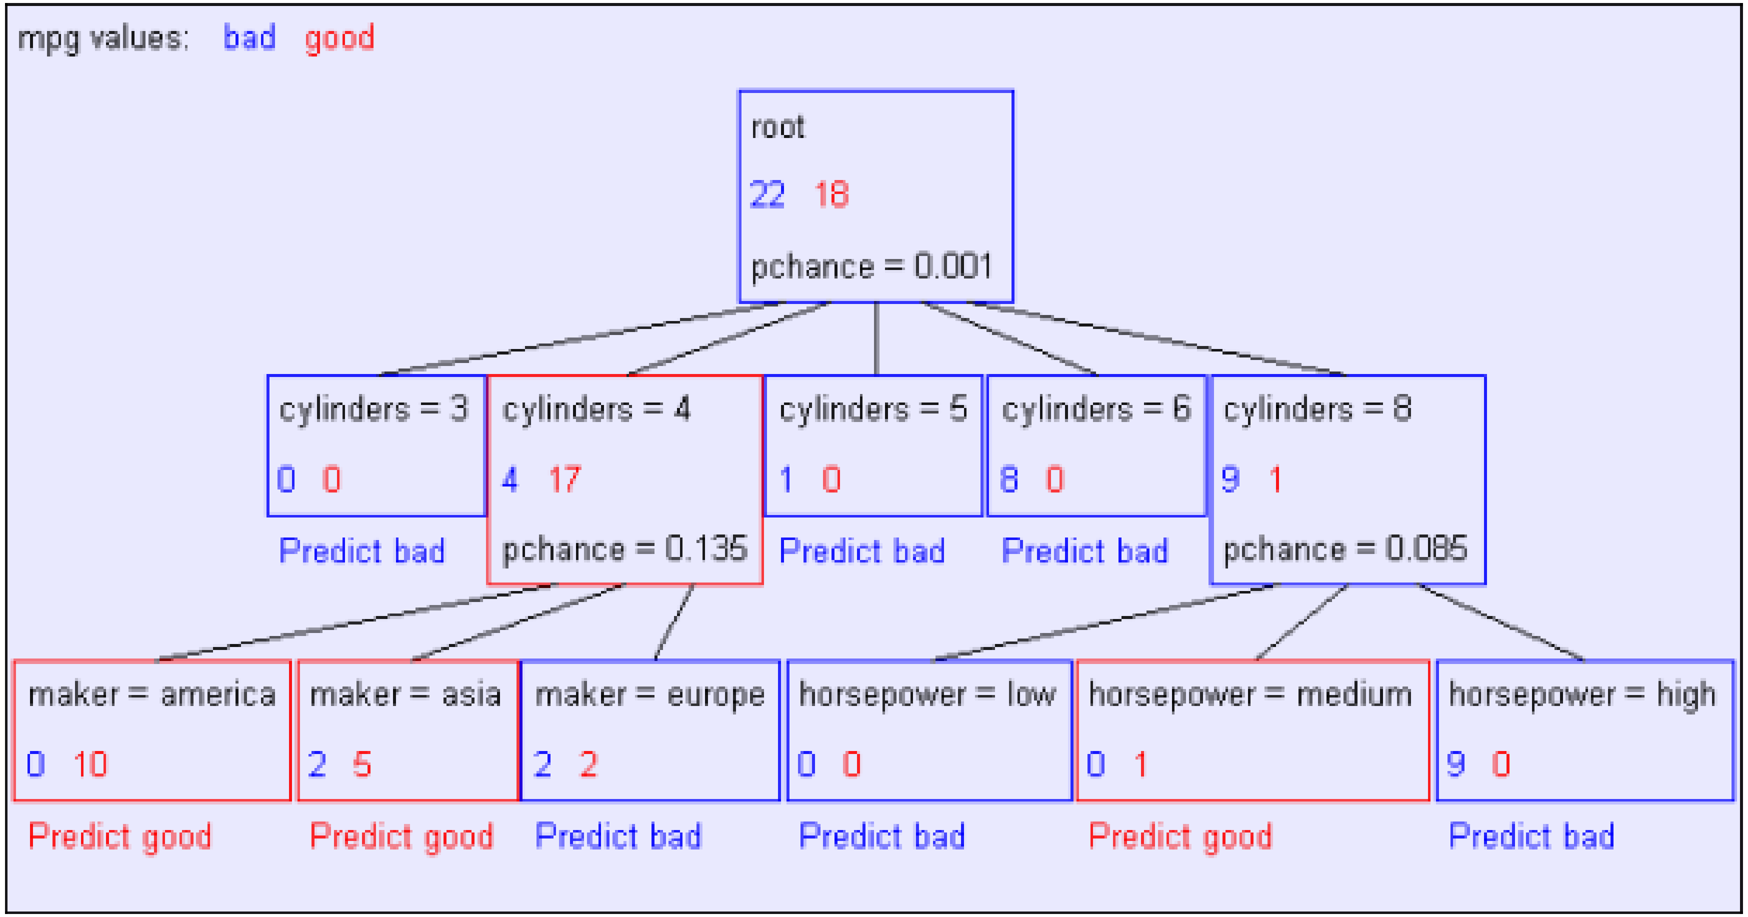

Final tree

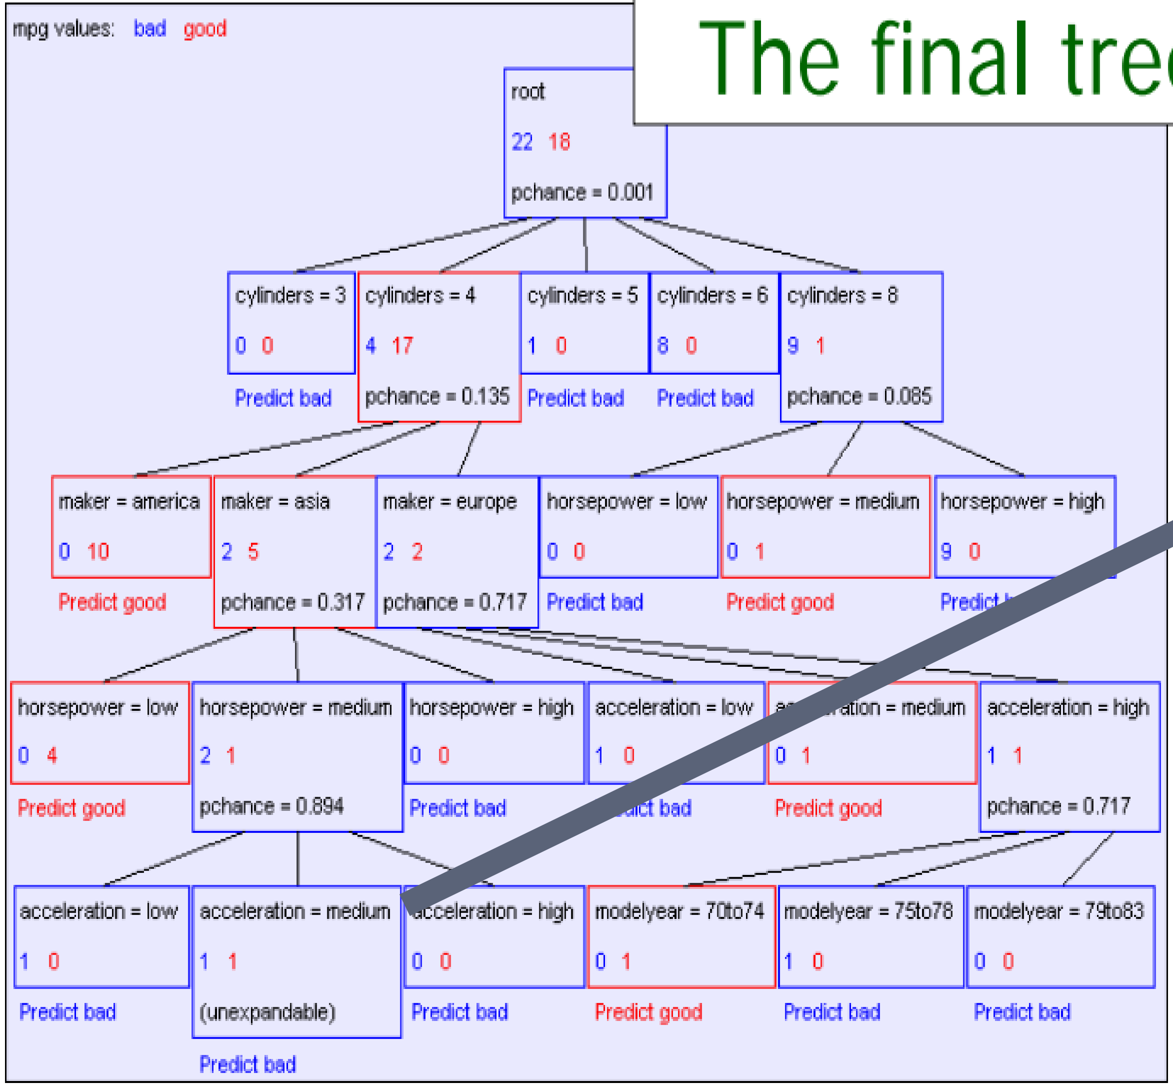

The final tree

Information gains using the training set (2 records)

mpg values: bad good

Input	Value	Distribution	Info Gain
cylinders	3		0
	4		
	5		
	6		
	8		
displacement	low		0
	medium		
	high		
horsepower	low		0
	medium		
	high		
weight	low		0
	medium		
	high		
acceleration	low		0
	medium		
	high		
modelyear	70to74		0
	75to78		
	79to83		
maker	america		0
	asia		
	europe		

Recap: Decision Tree algorithm

```
GrowTree(S):
    if y==0 for all <x,y> in S:
        return new leaf(0)
    else if y==1 for all <x,y> in S:
        return new leaf(1)
    else:
        x_j = max_info_gain(S)
        S0 = all <x,y> in S with x_j==0
        S1 = all <x,y> in S with x_j==1
        return new node(x_j, GrowTree(S0), GrowTree(S1))
```

Outline

- **Decision Trees**
  - Introduction
  - Representation
  - Algorithms
  - Splitting
  - Extended example
  - **Overfitting and Pruning**
  - Extensions

Overfitting

The final tree

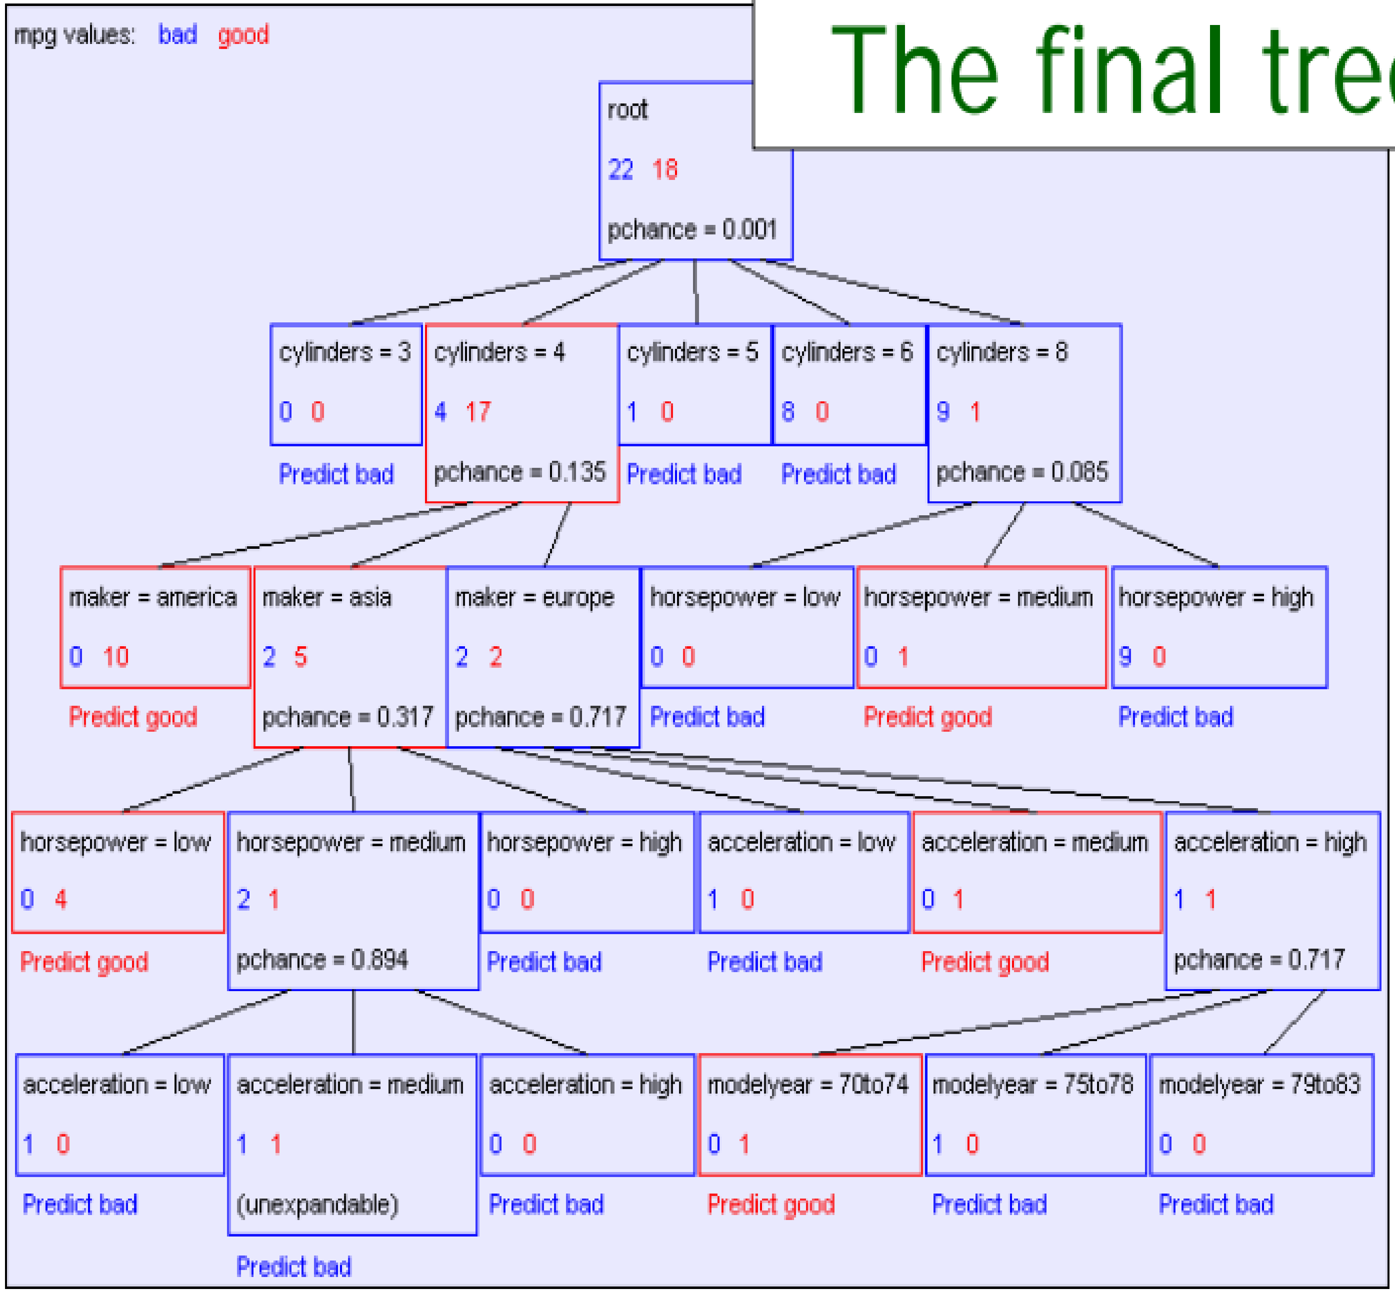

Overfitting

- Overfitting strikes again:

	Num Errors	Set Size	Percent Wrong
Training Set	1	40	2.50
Test Set	74	352	21.02

- How to deal with overfitting in Regression?
  - Regularization
- K-Nearest Neighbors?
  - Increasing K
- Naïve Bayes?
  - Smoothing

Overfitting in Decision Trees

- Three common solutions:
  1. Stop growing tree when split is not statistically significant
  2. Grow tree, then prune afterwards
  3. Set maximum depth

Example: Over-splitting

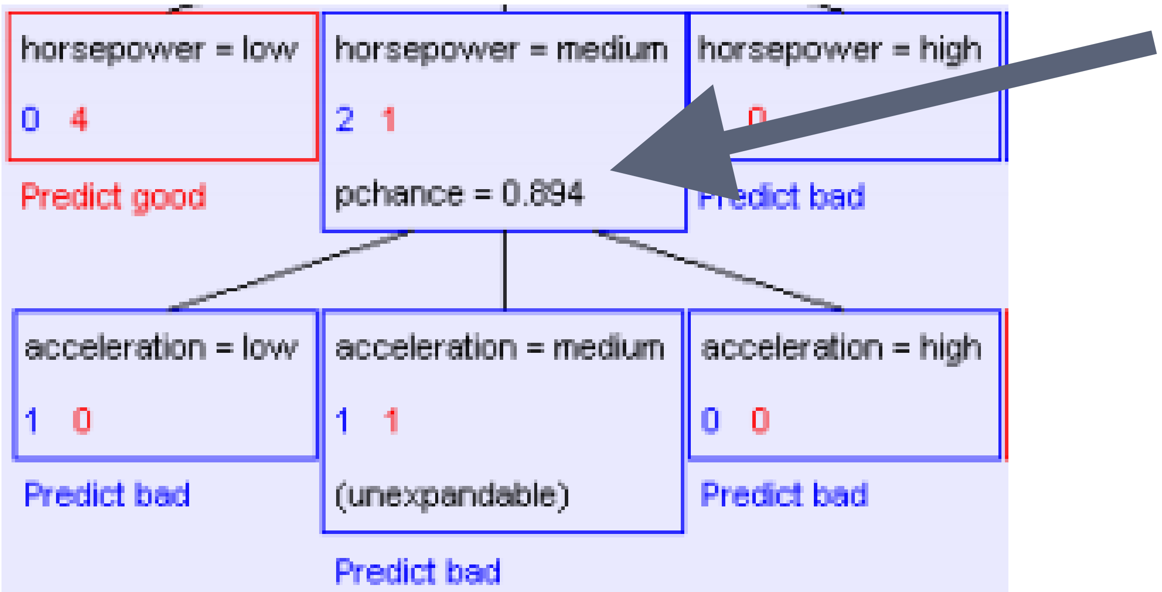

- Should we really split here?
- Only 3 relevant training examples
- The resulting distributions are likely due to chance

- **One solution**: compute the value/ significance of each split
  - For instance, a chi-squared test, or info gain
- Only split if value exceeds some threshold

$$\chi^2 = \sum_{i=1}^{n} \frac{(O_i - E_i)^2}{E_i}$$

Pruning

- Build the full decision tree
- Starting with the deepest nodes, delete splits where value of split does not exceed some threshold T
- Continue upward until no more prunable nodes

	Num Errors	Set Size	Percent Wrong
Training Set	1	40	2.50
Test Set	74	352	21.02

→

	Num Errors	Set Size	Percent Wrong
Training Set	5	40	12.50
Test Set	56	352	15.91

Regularization

- T is a regularization (hyper-)parameter
  - How to determine value?
- Cross-Validation!

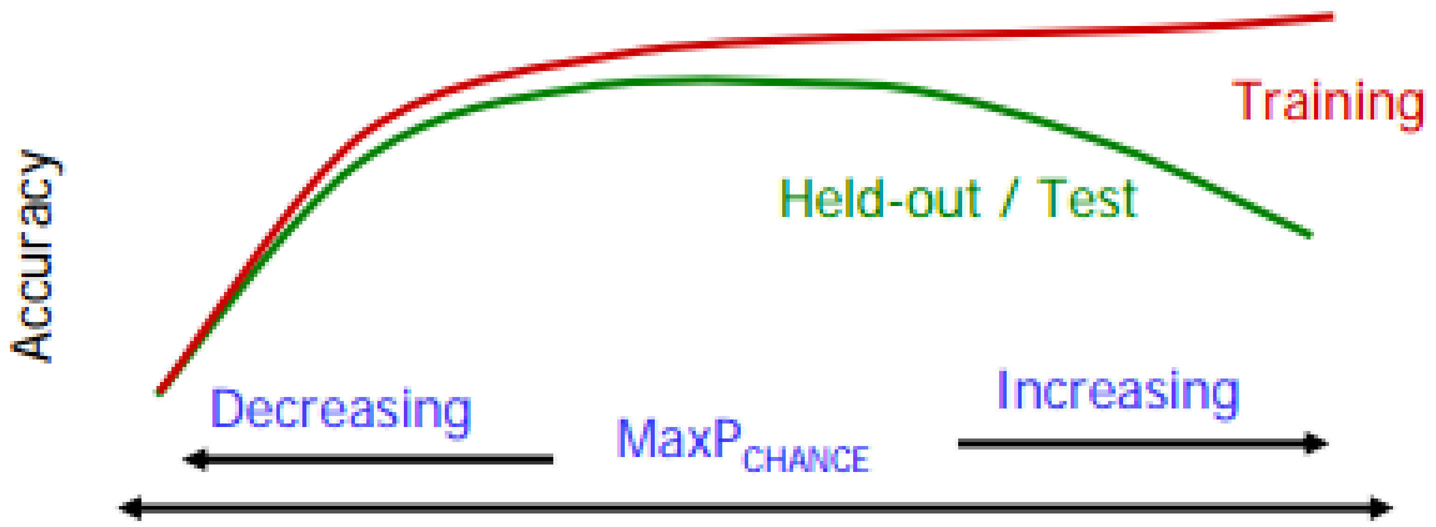

Outline

- Decision Trees
  - Introduction
  - Representation
  - Algorithms
  - Splitting
  - Overfitting and Pruning
  - **Extensions**

Multi-valued features

- Features with many discrete values:
  - Splits with many children (comparing Info Gain?)
  - Can produce degenerate cases
  - Common solution: One vs. all other values

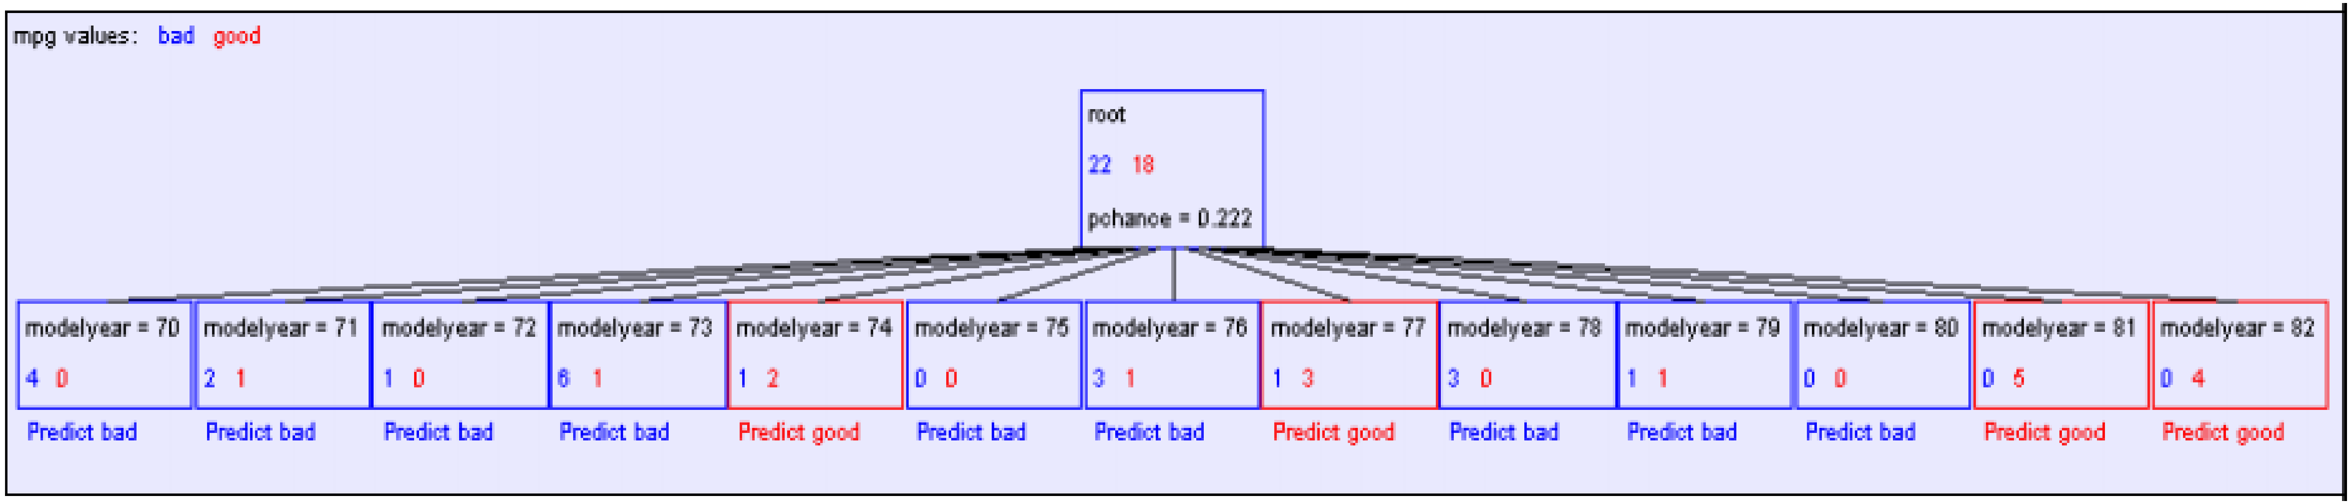

Continuous features

- Continuous features
  - Common solution: Bucket or threshold values
  - E.g., model years <1970, 1970-1980, >1980
- How to choose the buckets/thresholds?
  - Sort instances based on value of an attribute (e.g. year)
  - Identify adjacent examples that differ in their label
    - This generates a set of candidate thresholds splitting thresholds for that feature
  - Use information gain to decide appropriate threshold

Thresholded splits

- Bucketing example:
  - Creates deeper, denser tree (for same value of T)

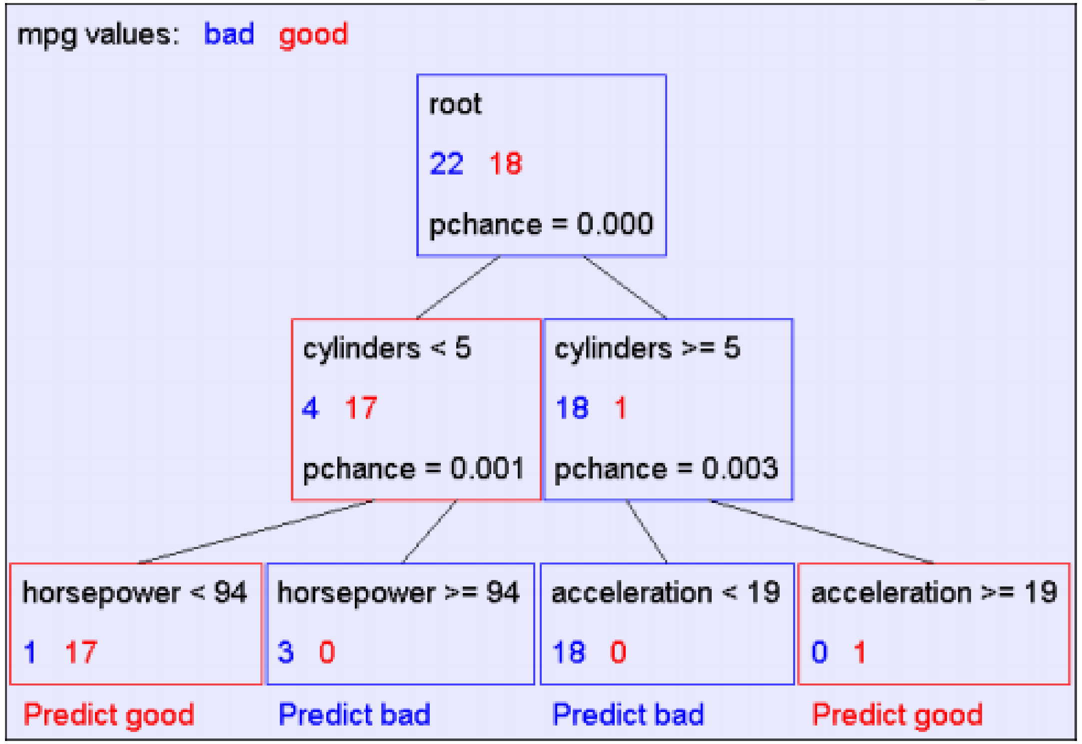

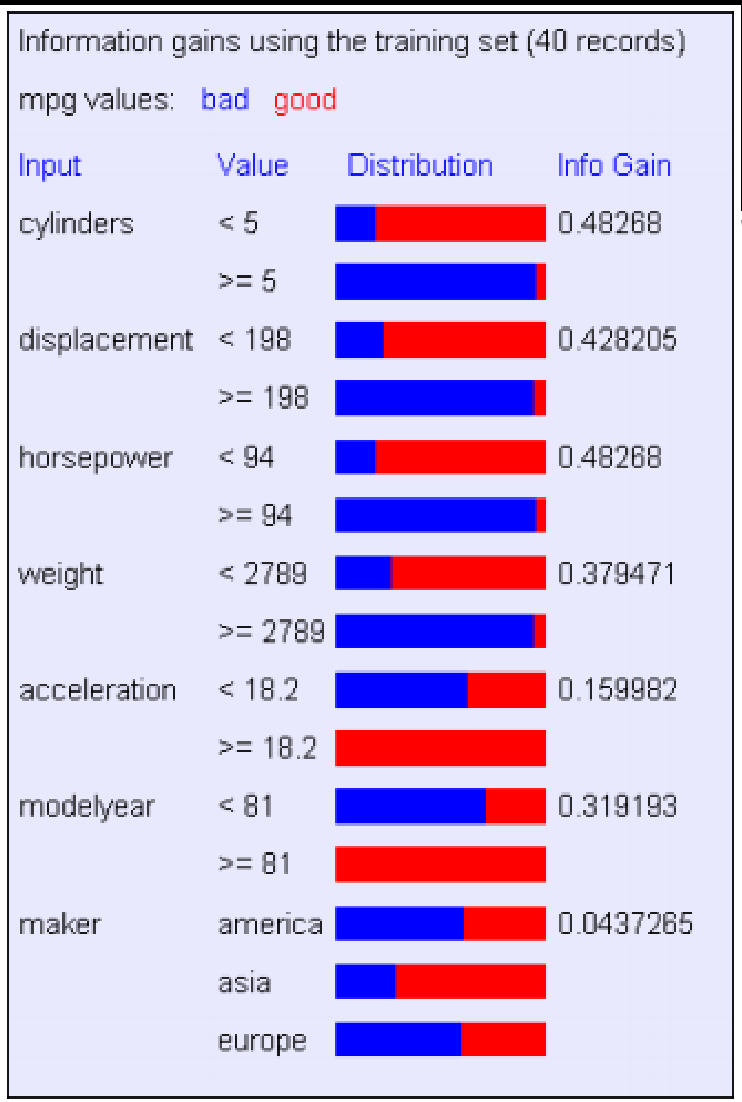

	Num Errors	Set Size	Percent Wrong
Training Set	1	40	2.50
Test Set	53	352	15.06

Output probabilities

- How to do better than predicting majority class?
- Estimate probabilities from the relevant examples at each node
- Can use smoothing to improve estimates (e.g., Laplace smoothing)

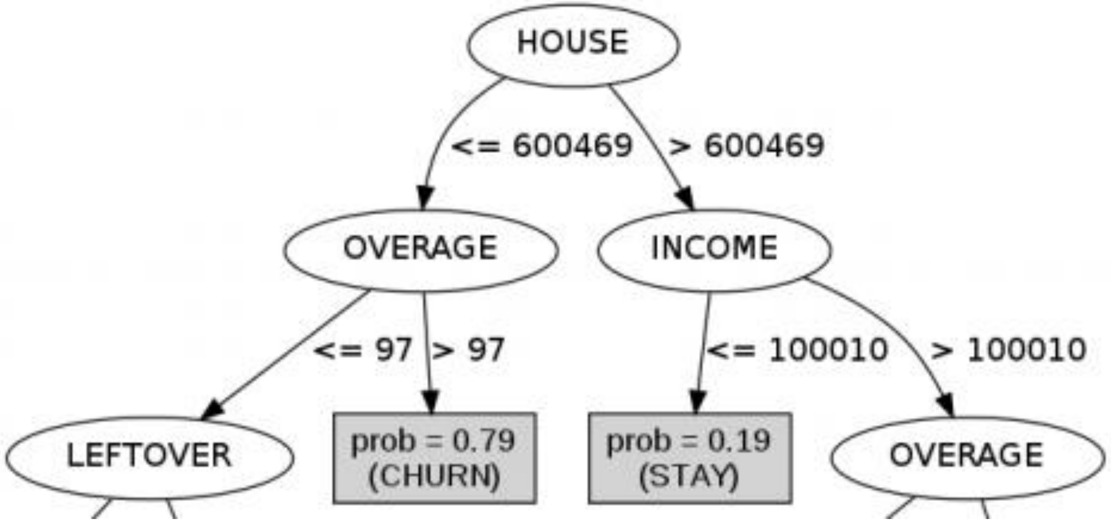

Scaling up

- More data is almost always better

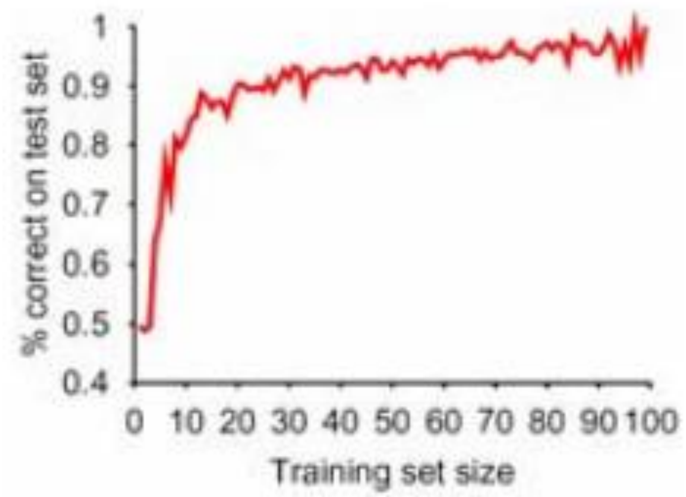

- Scaling up with standard recursive algorithms can be hard
- New algorithms make single pass through data
  - E.g. Very Fast Decision Trees (VFTD)