



THAT WAS SURPRISINGLY EASY. HOW COME THE ROBOTIC UPRISING USED SPEARS AND ROCKS INSTEAD OF MISSILES AND LASERS?

IF YOU LOOK TO HISTORICAL DATA, THE VAST MAJORITY OF BATTLE-WINNERS USED PRE-MODERN WEAPONRY.

Thanks to machine-learning algorithms, the robot apocalypse was short-lived.

INFO 251: Applied Machine Learning

Trees and Forests

Key Concepts (last lecture)

- Churn prediction
- Decision boundaries
- Hyper-rectangles
- Splitting
- Information gain
- Recursive tree building
- Overfitting trees
- Pruning trees

Course Outline

- Causal Inference and Research Design
 - Experimental methods
 - Non-experiment methods
- Machine Learning
 - Design of Machine Learning Experiments
 - Linear Models and Gradient Descent
 - **Non-linear models**
 - Neural models
 - Unsupervised Learning
 - Practicalities, Fairness, Bias
- Special topics

Outline

- Decision Trees: Loose ends
 - **Overfitting and Pruning**
 - Extensions
- Regression Trees
- Random Forests
- Feature Importance

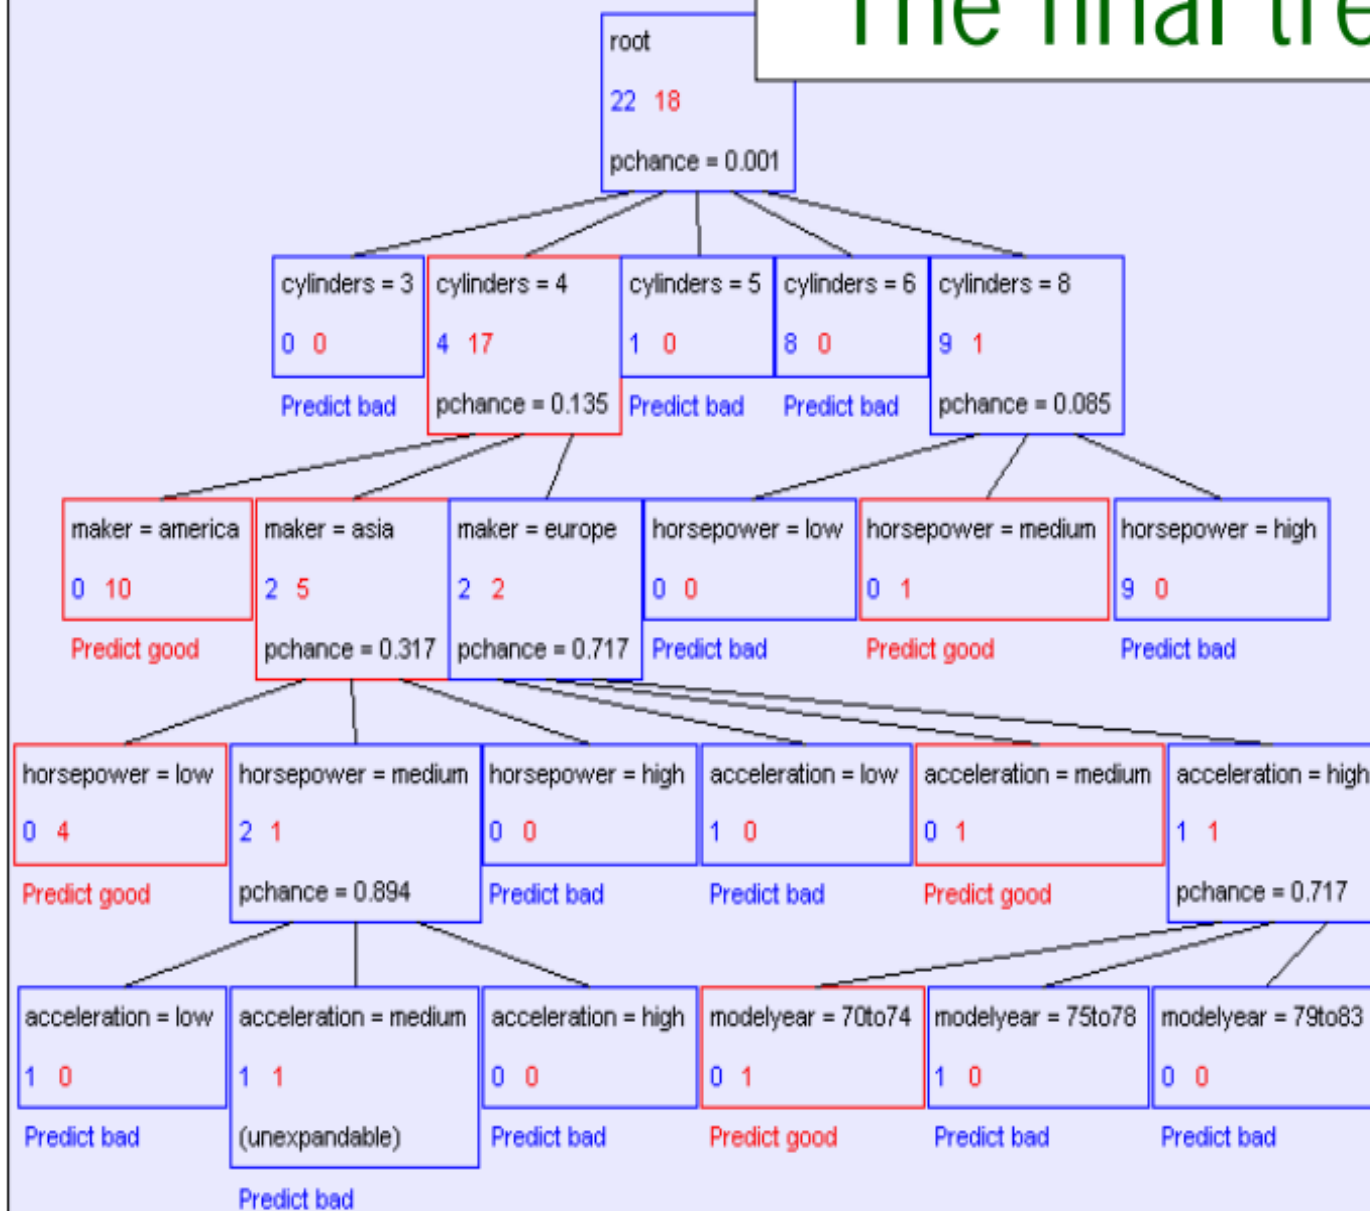
Recap: Decision Tree algorithm

```
GrowTree(S) :  
    if y==0 for all <x,y> in S:  
        return new leaf(0)  
    else if y==1 for all <x,y> in S:  
        return new leaf(1)  
    else:  
         $x_j = \text{max\_info\_gain}(S)$   
        S0 = all <x,y> in S with  $x_j == 0$   
        S1 = all <x,y> in S with  $x_j == 1$   
        return new node( $x_j$ , GrowTree(S0), GrowTree(S1))
```

Overfitting

The final tree

mpg values: bad good



Overfitting

- Overfitting strikes again:

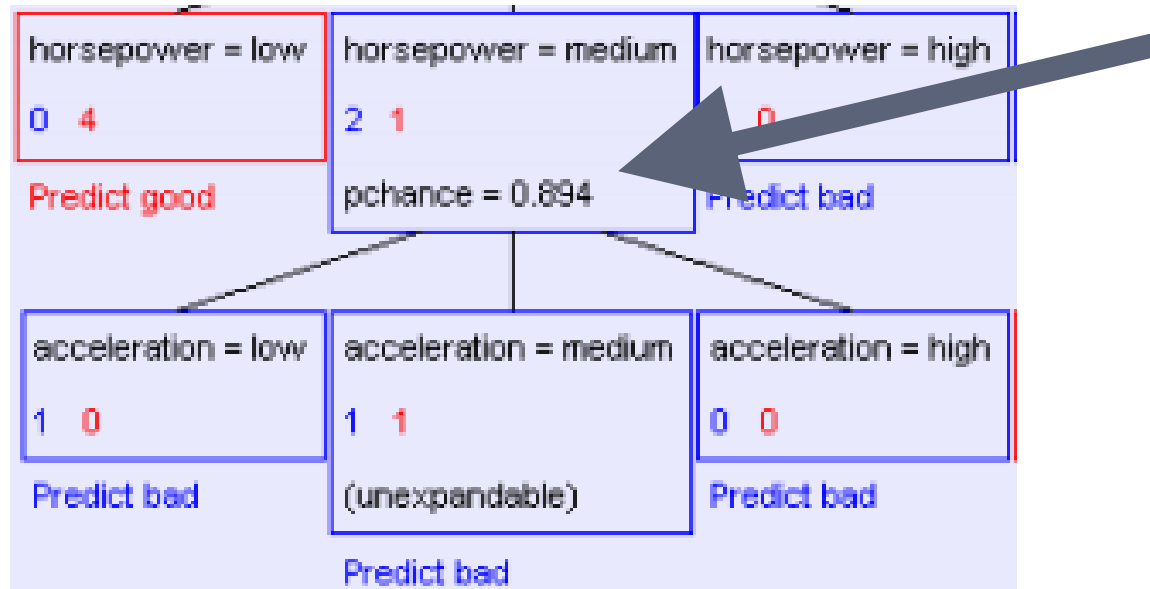
	Num Errors	Set Size	Percent Wrong
Training Set	1	40	2.50
Test Set	74	352	21.02

- How to deal with overfitting in Regression?
 - Regularization
- K-Nearest Neighbors?
 - Increasing K
- Naïve Bayes?
 - Smoothing

Overfitting in Decision Trees

- Three common solutions:
 1. Stop growing tree when split is not statistically significant
 2. Grow tree, then prune afterwards
 3. Set maximum depth

Example: Over-splitting



- Should we really split here?
- Only 3 relevant training examples
- The resulting distributions are likely due to chance

- **One solution:** compute the value/ significance of each split
 - For instance, a chi-squared test, or info gain
- Only split if value exceeds some threshold

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}$$

Pruning

- Build the full decision tree
- Starting with the deepest nodes, delete splits where value of split does not exceed some threshold T
- Continue upward until no more prunable nodes

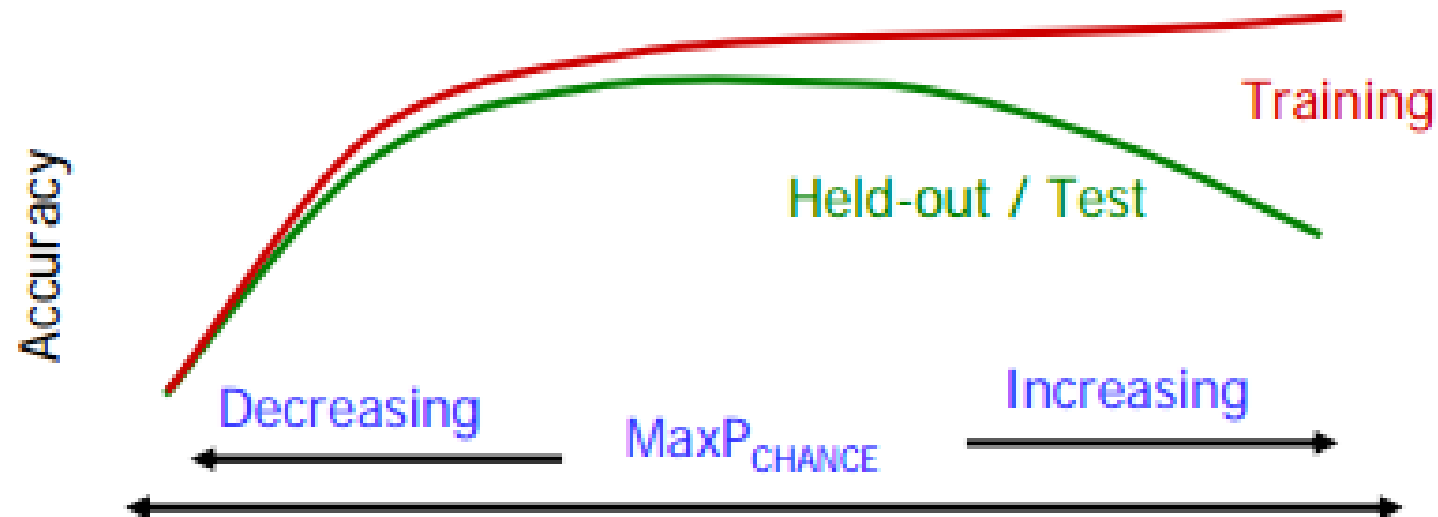
		Num Errors	Set Size	Percent Wrong
Training Set	1	40		2.50
Test Set	74	352		21.02



		Num Errors	Set Size	Percent Wrong
Training Set	5	40		12.50
Test Set	56	352		15.91

Regularization

- T is a regularization (hyper-)parameter
 - How to determine value?
- Cross-Validation!

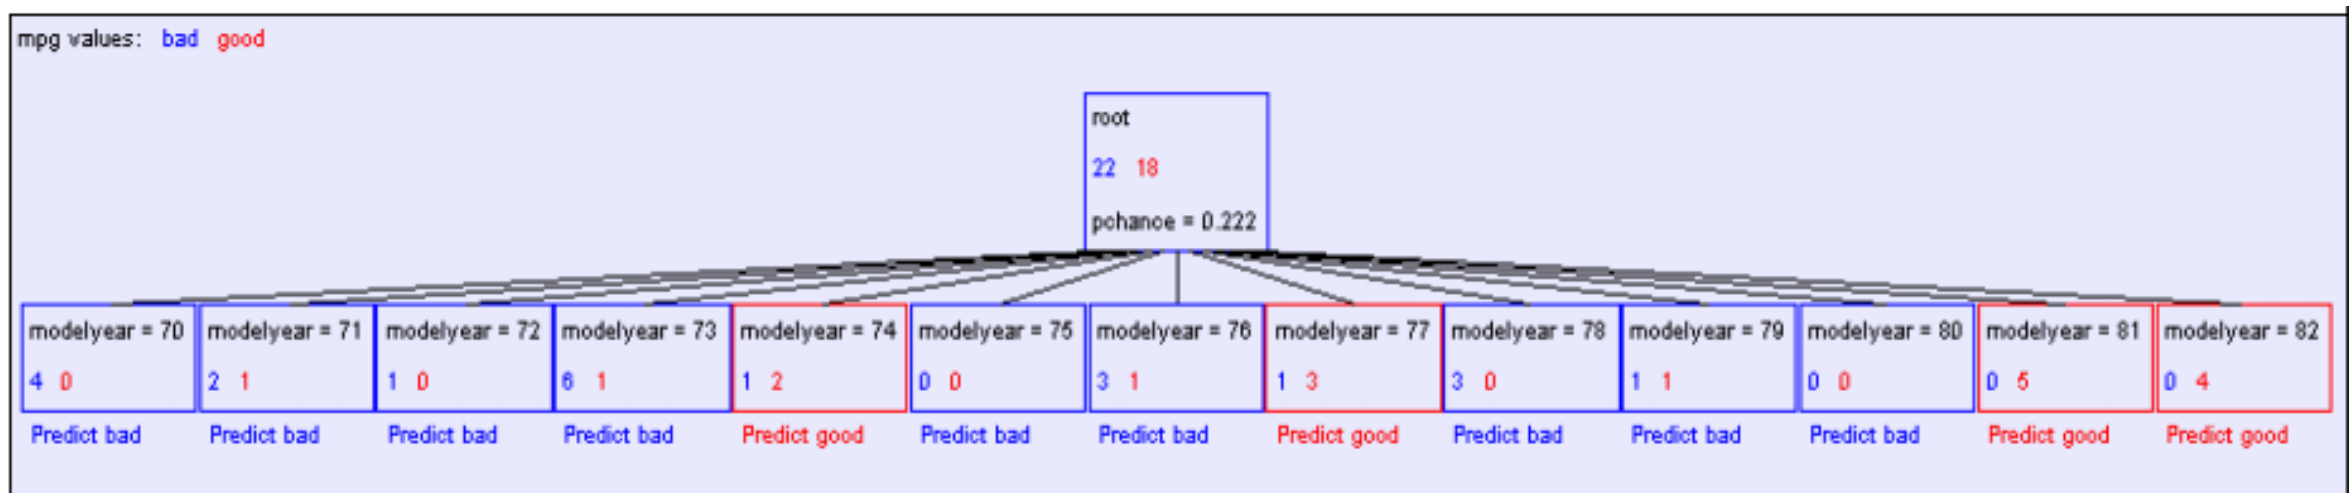


Outline

- Decision Trees: Loose ends
 - Overfitting and Pruning
 - **Extensions**
- Regression Trees
- Random Forests
- Feature Importance

Multi-valued features

- Features with many discrete values:
 - Splits with many children (comparing Info Gain?)
 - Can produce degenerate cases
 - Common solution: One vs. all other values

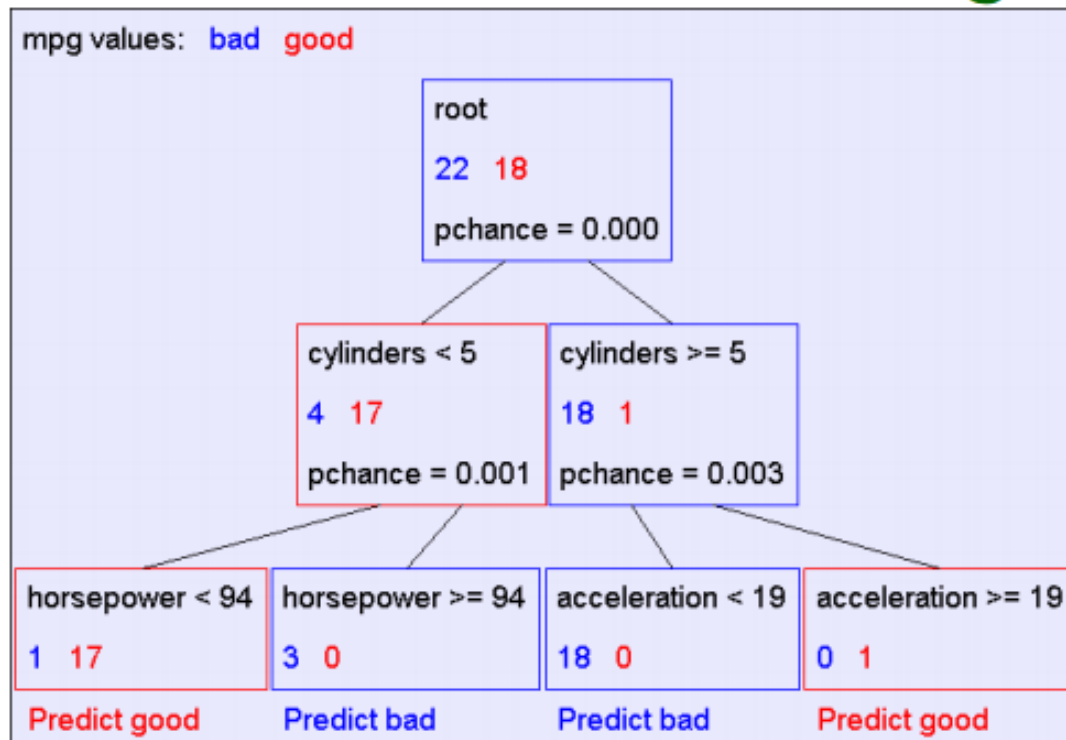


Continuous features

- Continuous features
 - Common solution: Bucket or threshold values
 - E.g., model years <1970, 1970-1980, >1980
- How to choose the buckets/thresholds?
 - Sort instances based on value of an attribute (e.g. year)
 - Identify adjacent examples that differ in their label
 - This generates a set of candidate thresholds splitting thresholds for that feature
 - Use information gain to decide appropriate threshold

Thresholded splits

- Bucketing example:
 - Creates deeper, denser tree (for same value of T)



Information gains using the training set (40 records)

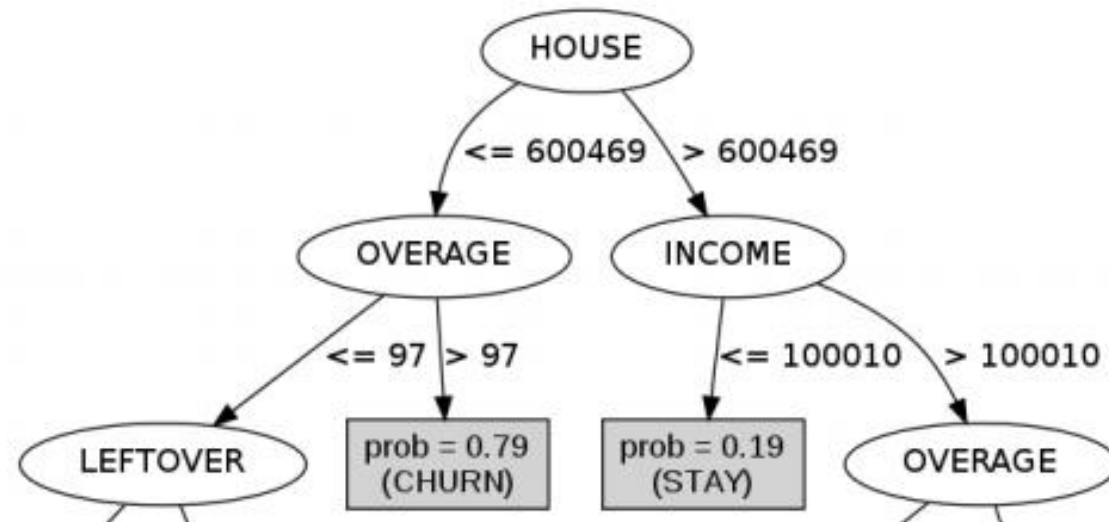
mpg values: bad good

Input	Value	Distribution	Info Gain
cylinders	< 5		0.48268
	>= 5		
displacement	< 198		0.428205
	>= 198		
horsepower	< 94		0.48268
	>= 94		
weight	< 2789		0.379471
	>= 2789		
acceleration	< 18.2		0.159982
	>= 18.2		
modelyear	< 81		0.319193
	>= 81		
maker	america		0.0437265
	asia		
	europa		

		Num Errors	Set Size	Percent Wrong
Training Set	1	40		2.50
Test Set	53	352		15.06

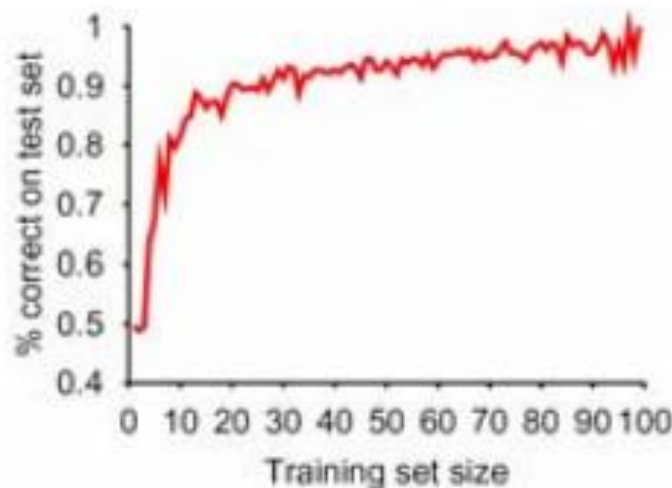
Output probabilities

- How to do better than predicting majority class?
- Estimate probabilities from the relevant examples at each node
- Can use smoothing to improve estimates (e.g., Laplace smoothing)



Scaling up

- More data is almost always better



- Scaling up with standard recursive algorithms can be hard
- New algorithms make single pass through data
 - E.g. Very Fast Decision Trees (VFTD)

Outline

- Decision Trees: Loose ends
 - Overfitting and Pruning
 - Extensions
- **Regression Trees**
- Random Forests
- Feature Importance

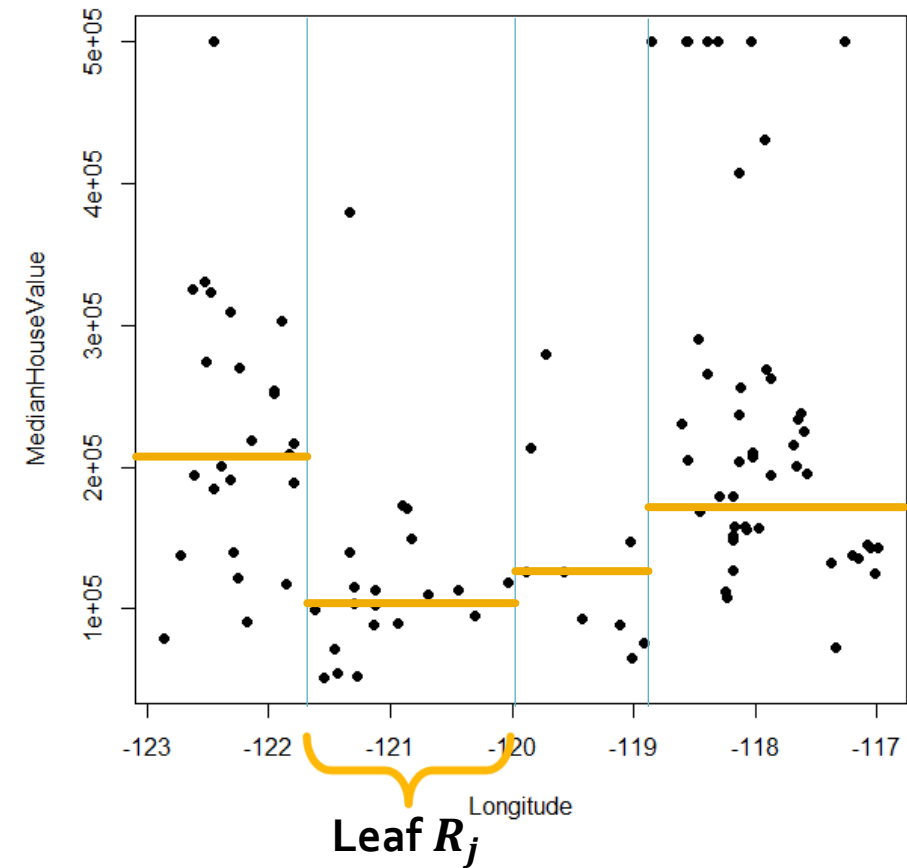
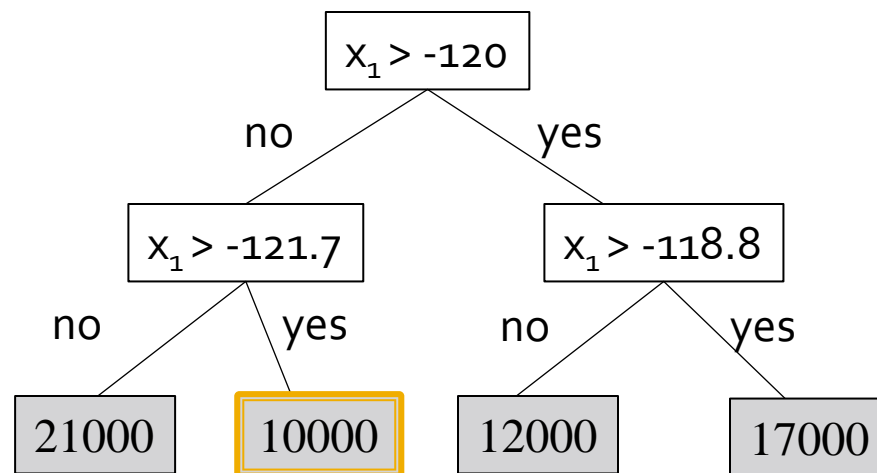
Regression Trees

- What if output values are continuous or real-valued (i.e., not discrete)?
- Regression trees
 - Construct binary tree, minimize error in each leaf
 - Before, we counted # elements of each type in leaf
 - Now we choose predicted value that minimizes error
- Example: Predict median housing value based on a house's location (latitude, longitude)

Regression Trees

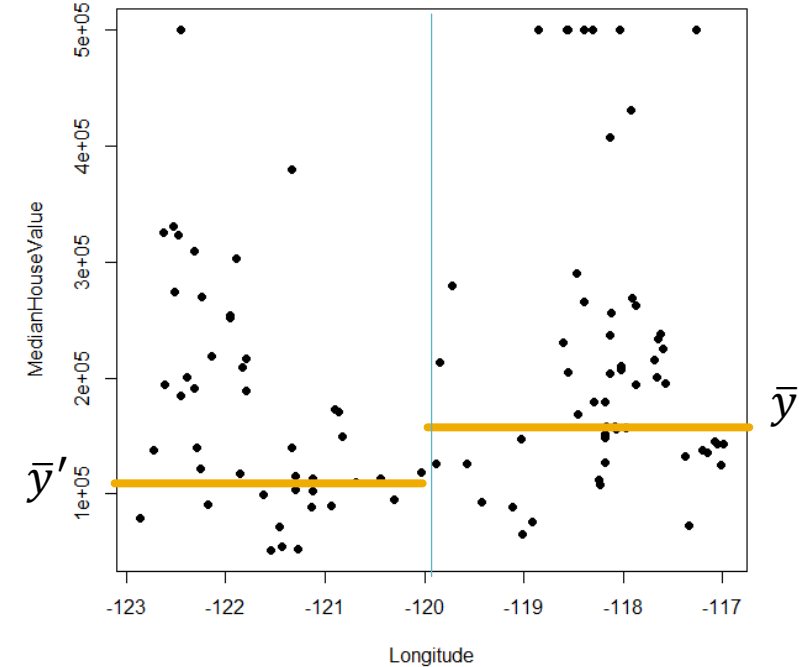
```
> head(calif[,c(1,8,9)])
```

	MedianHouseValue	Latitude	Longitude
1	452600	37.88	-122.23
2	358500	37.86	-122.22
3	352100	37.85	-122.24
4	341300	37.85	-122.25
5	342200	37.85	-122.25
6	269700	37.85	-122.25



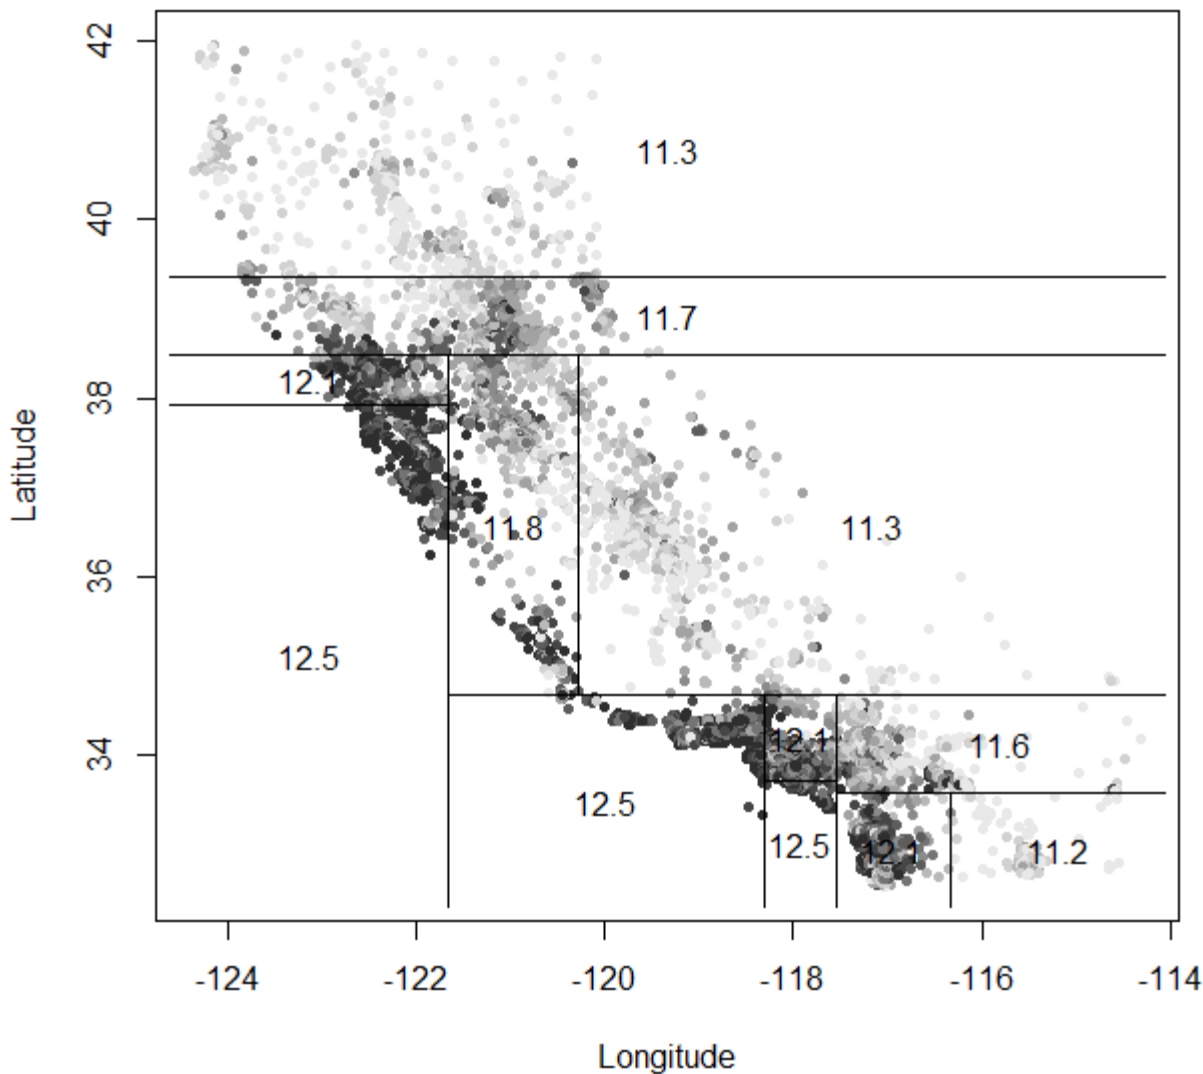
Regression Trees

- How to choose split point?
 - Idea: Minimize prediction error
- In 1-dimension: choose s to minimize
 - $\min_{\bar{y}} \sum_{i: x_i > s} (\bar{y} - y_i)^2 + \min_{\bar{y}'} \sum_{i: x_i \leq s} (\bar{y}' - y_i)^2$
 - Consider finite splits (e.g. s between data)
- This intuition generalizes trivially to D dimensions



Regression Trees: Recursive Algorithm

1. Start with a single node (c_j) containing all points.
 1. Calculate predicted value $\bar{y}_{c_j} = \frac{1}{n} \sum_{i \in c_j} y_i$
 2. Calculate total error: $J = \sum_{c_j} \sum_{i \in c_j} (\bar{y}_{c_j} - y_i)^2$
2. If all points in the node have the identical features, stop.
 - Otherwise, search all binary splits of all variables for split that most reduces J
 - Stop if J decreases less than δ or if nodes are close to empty
 - Otherwise, make that split, creating two new nodes
3. Recurse on each new node

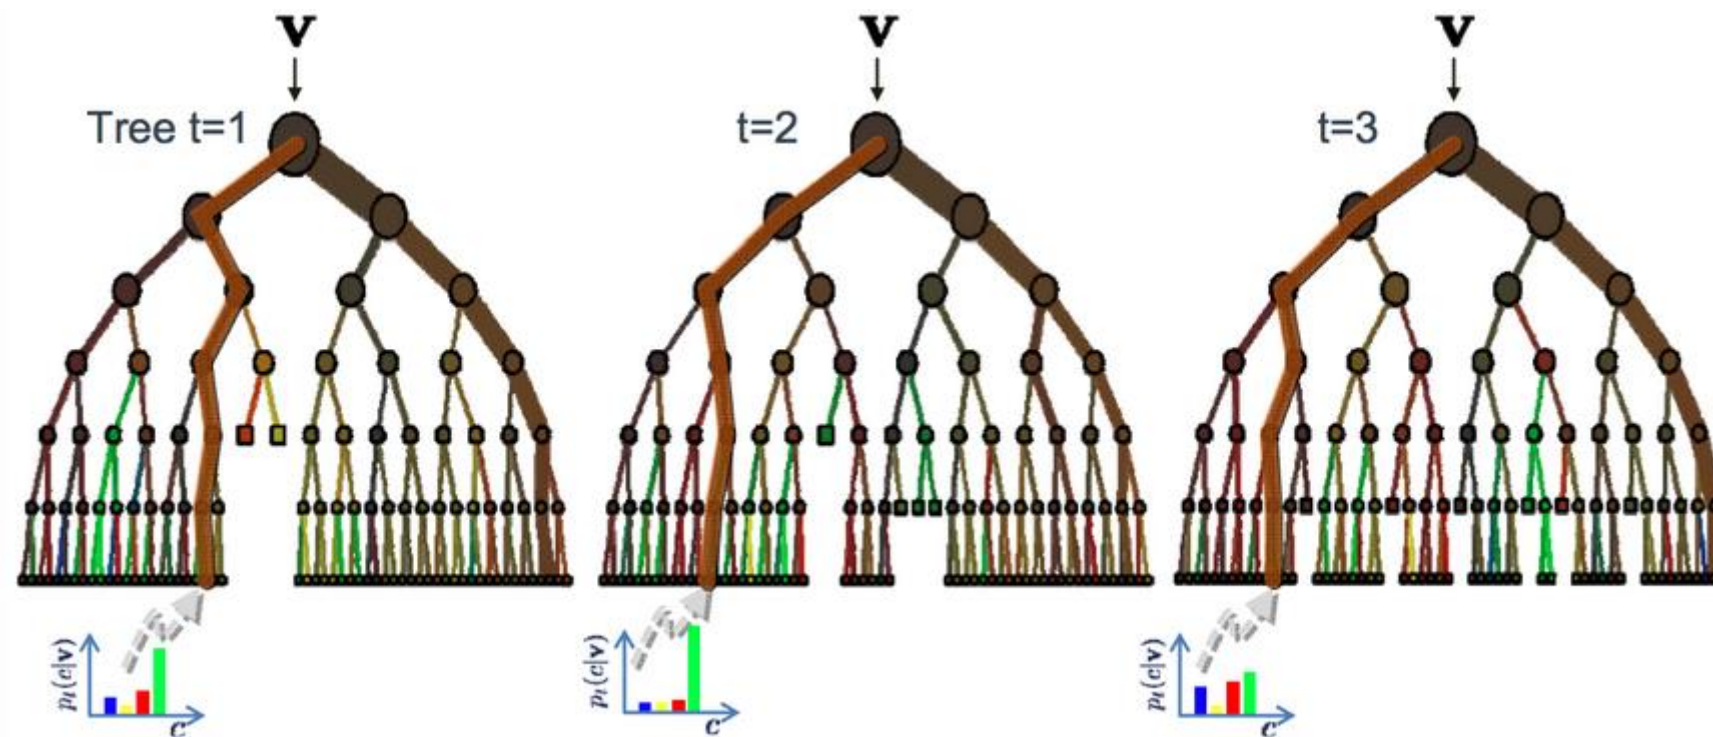


Outline

- Decision Trees: Loose ends
 - Overfitting and Pruning
 - Extensions
- Regression Trees
- **Random Forests**
- Feature Importance

Trees to forests

- Which classifier works best?
 - “Random forests” combine outputs of multiple classifiers



Building a forest

- Bootstrap sample a new training set
 - with replacement
- Build a decision tree
 - Random subset of splits/features
 - Forces differentiated trees
 - No pruning! “Regularization” through forest
- Repeat until you have lots of trees
- Predict by taking a vote among the trees

Example: The “CART” forest

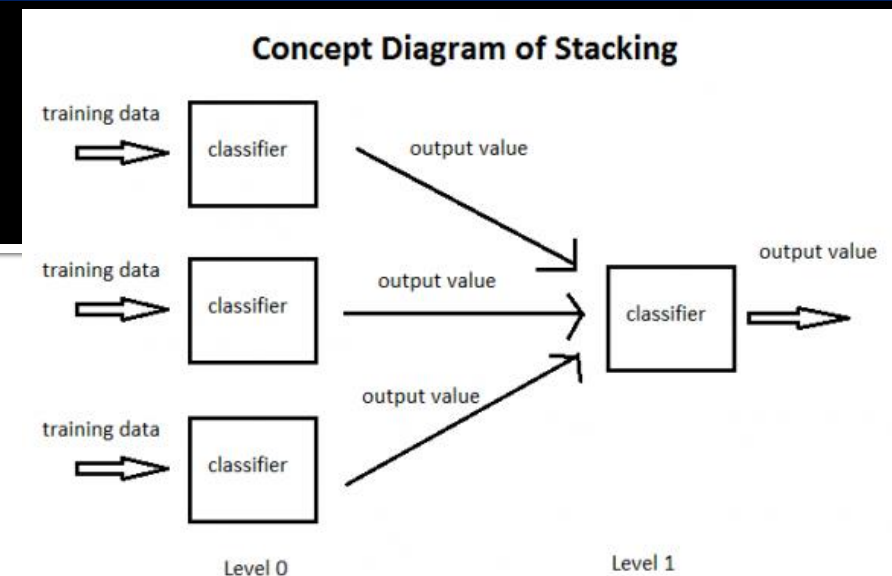
- Formally:

$$\hat{y}_i = \sum_{f_k \in \mathcal{F}} f_k(x_i)$$

- \mathcal{F} is the space of regression trees
 - Each f_k maps data examples x_i to tree leaves
 - With CART (a common decision tree algorithm), each leaf isn't a decision value, it's a “score”
 - Scores are summed across trees

Ensemble methods

- Bagging = **bootstrap aggregating**
 - Create artificial versions of data via bootstrap
 - 1 sample = bootstrap
 - M samples = bagging
- Stacking: train model (e.g. another tree, a logistic regression) on output of other models
- Boosting: Can a set of weak learners create a single strong learner? (Kearns, 1988)
 - Train a sequence of models, each emphasizes the examples misclassified by the previous model



Adaptive Boosting (AdaBoost)

- Adaboost:

1. Initially, set a weight for each training example = $1/n$
2. Train a classifier where the objective respects the weights
3. Increase the weights for misclassified examples
4. Return to 2

Gradient Boosting



Linear Regression



Gradient Boosting

<http://blog.kaggle.com/2017/01/23/a-kaggle-master-explains-gradient-boosting/>

(Extreme) Gradient Boosting

- Start with regression tree-based model:

$$\hat{y}_i = \sum_{f_k \in \mathcal{F}} f_k(x_i)$$

- Gradient boosting loss function “fits on residuals”:

$$\mathcal{L}^{(t)} = \sum_{i=1}^n l(y_i, \underbrace{\hat{y}_i^{(t-1)} + f_t(x_i)}_{f_t \text{ fits on residual of } t-1}) + \underbrace{\gamma T + \frac{1}{2} \lambda \|w\|^2}_{\text{Regularization penalty}}$$

- T is the number of leaves
 - t indexes training iterations
 - w is vector of scores on each leaf (i.e., the leaf weights)
- Optimization is similar to gradient descent
 - Relies on being able to measure how good each tree is
 - Next tree solves for the loss of prior tree

Outline

- Decision Trees: Loose ends
 - Overfitting and Pruning
 - Extensions
- Regression Trees
- Random Forests
- **Feature Importance**

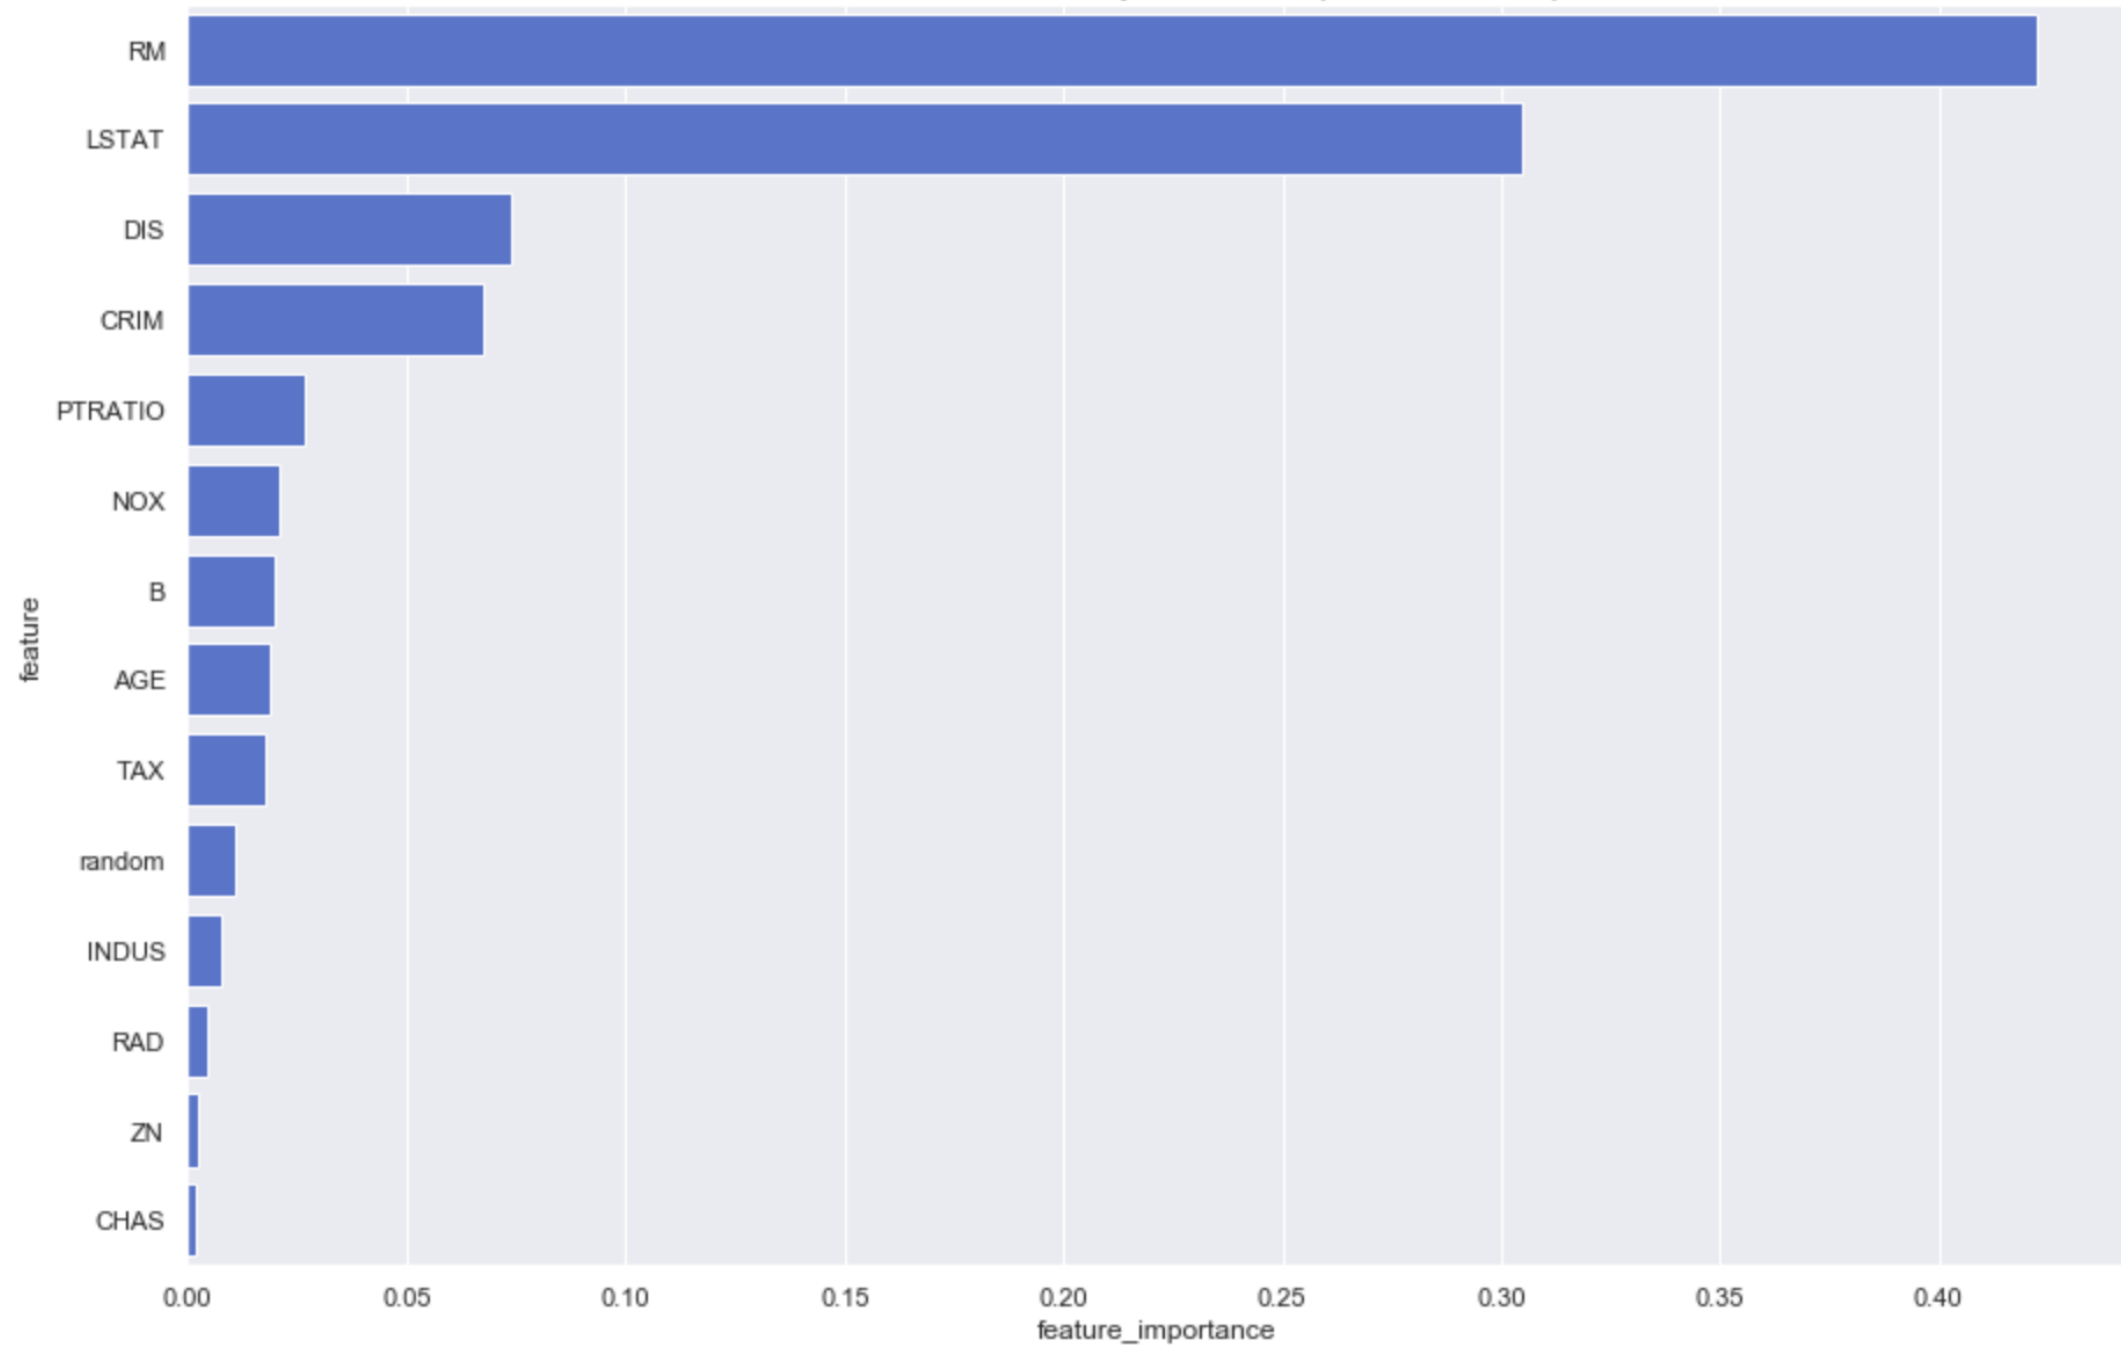
Feature Importance

- The primary focus of random forests is prediction, rather than inference
 - A single tree is fairly interpretable, but it's hard to interpret a forest
 - This is generally true for complex, non-parametric, and non-linear models
- Nonetheless, people frequently still want to do some ex-post interpretation
 - "What features are important to the classifier?"

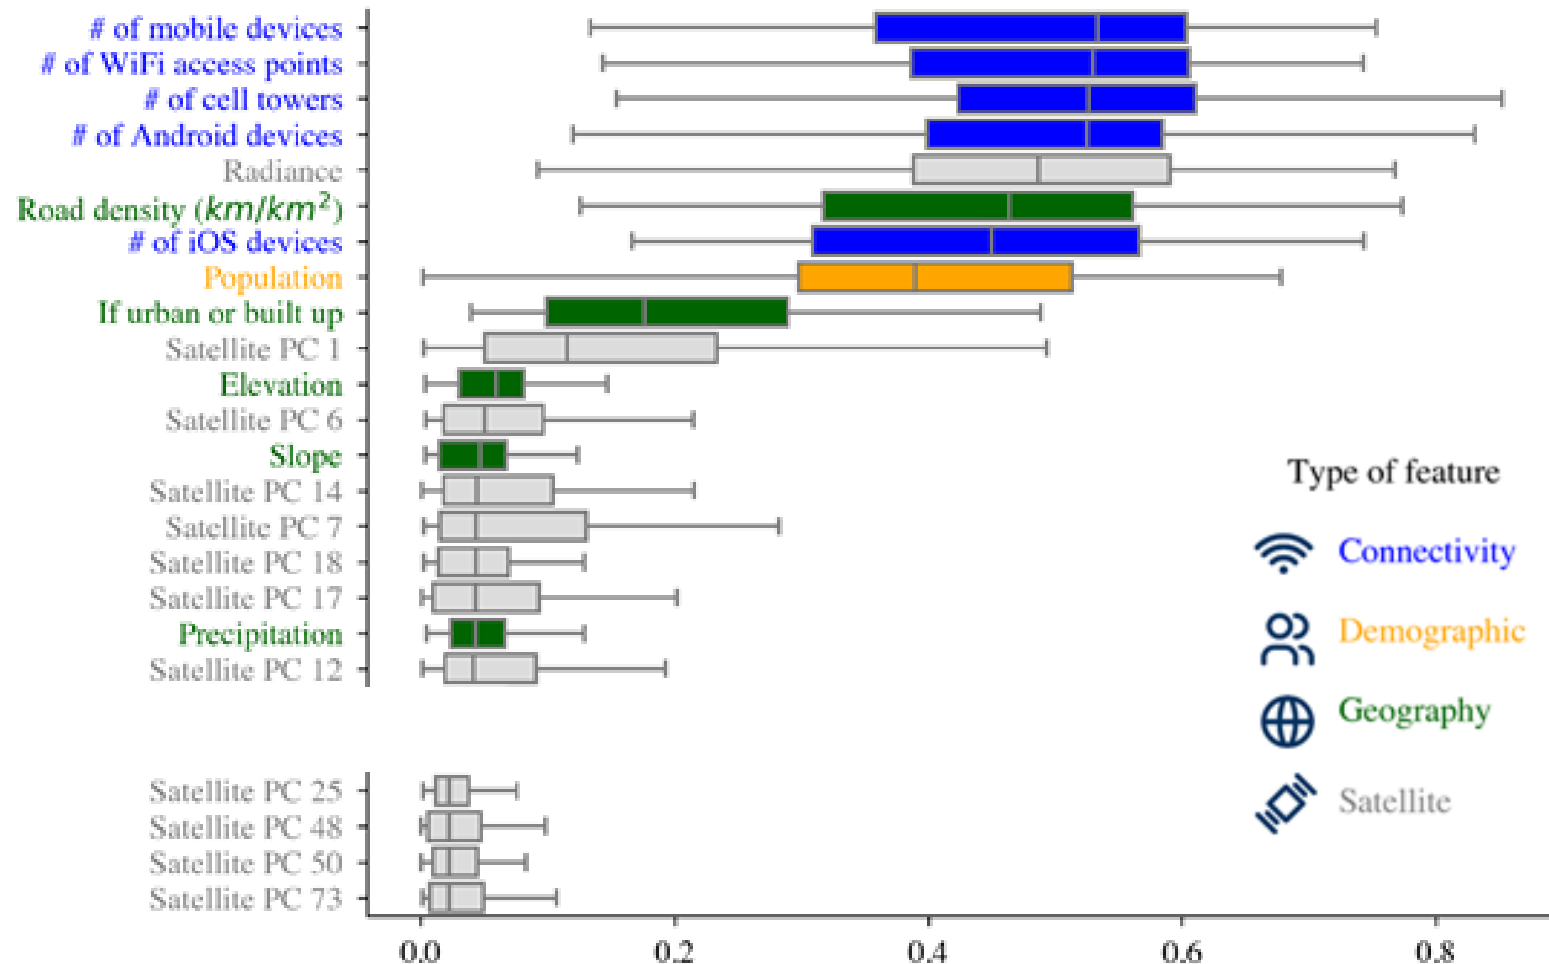
Feature Importance

- Intuitively, the features that “matter”:
 - Occur high in the tree (high information gain for that tree)
 - Occur frequently in the tree (if feature is non-binary)
 - Occur in many trees (if it's a forest)

Default feature importance (scikit-learn)



Feature Importance: Example



Feature Importance

- Formally, two common approaches:
 1. **Mean decrease impurity (aka Gini importance):** average (across trees) decrease in weighted impurity caused by that feature

```

1  from sklearn.datasets import load_boston
2  from sklearn.ensemble import RandomForestRegressor
3  import numpy as np
4  #Load boston housing dataset as an example
5  boston = load_boston()
6  X = boston["data"]
7  Y = boston["target"]
8  names = boston["feature_names"]
9  rf = RandomForestRegressor()
10 rf.fit(X, Y)
11 print "Features sorted by their score:"
12 print sorted(zip(map(lambda x: round(x, 4), rf.feature_importances_), names),
13               reverse=True)

```

Features sorted by their score:

```

[(0.5298, 'LSTAT'), (0.4116, 'RM'), (0.0252, 'DIS'), (0.0172, 'CRIM'), (0.0065, 'NOX'),
(0.0035, 'PTRATIO'), (0.0021, 'TAX'), (0.0017, 'AGE'), (0.0012, 'B'), (0.0008, 'INDUS'),
(0.0004, 'RAD'), (0.0001, 'CHAS'), (0.0, 'ZN')]

```

Feature Importance

- Issues with impurity:
 - Biased towards features with multiple values
 - What happens when two features are closely correlated?

```

1  size = 10000
2  np.random.seed(seed=10)
3  X_seed = np.random.normal(0, 1, size)
4  X0 = X_seed + np.random.normal(0, .1, size)
5  X1 = X_seed + np.random.normal(0, .1, size)
6  X2 = X_seed + np.random.normal(0, .1, size)
7  X = np.array([X0, X1, X2]).T
8  Y = X0 + X1 + X2
9
10 rf = RandomForestRegressor(n_estimators=20, max_features=2)
11 rf.fit(X, Y);
12 print "Scores for X0, X1, X2:", map(lambda x:round (x,3),
13                                     rf.feature_importances_)

```

Scores for X0, X1, X2: [0.278, 0.66, 0.062]

Feature Importance

- Formally, two common approaches:
 1. **Mean decrease impurity**: average (across trees) decrease in weighted impurity caused by that feature
 2. **Mean decrease accuracy (“Permutation Importance”)**: average (across trees) decrease in performance when a given feature is randomized
 - Not implemented in sklearn (but very easy to do so by hand - see ESLII reading)

Feature Importance

- Issues
 - Interpret feature importances at your own risk!
 - They are informative, but rather atheoretical

Key Concepts (this lecture)

- Overfitting and pruning
- Regression trees
- Random forests
- AdaBoost
- Gradient boosting
- Feature Importance

Recap

- Regression
 - Parametric, fast training, linear
- Nearest Neighbors
 - Non-parametric, no training, complex decisions
- Naïve Bayes
 - Parametric, very fast training
- Decision Trees
 - Non-linear decisions, intuitive model

For Next Class:

- Read:
 - Daume, chapters 4 and 10
- Keep working on Problem Set 4!

